



KTH Information and
Communication Technology

Label-Efficient Multi-Objective Machine Learning for e-Commerce

Exploration of transfer, multi-objective and active learning
across shallow and deep neural network architectures
for e-commerce product classification.

MATTIAS ARRO

Master's Thesis at KTH Information and Communication Technology
MSc Data Science (EIT Digital track)

Academic Examiner: Magnus Boman
Academic Supervisor: Jim Dowling
Industrial Supervisor: Abubakrelsedik Karali

2018

Abstract

Neural networks are powerful and flexible models capable of highly accurate predictions, transfer and multi-objective learning, though their weakness is the amount of labelled data needed to train them. This work looks at three ways to tackle this label complexity problem: semi-supervised, transfer and active learning. The former is described theoretically, and the usefulness of the latter two is evaluated through a series of experiments. First, a combination of deep and linear models with different input features and representations thereof is trained on a retail product classification problem. Transfer learning is evaluated (1) subjectively by using pre-trained 2D CNNs to extract dense embeddings from images, to be used for calculating visual similarity scores, and (2) by using pre-trained 2D CNNs and deep averaging networks for the product classification task. Multi-objective learning is evaluated on a series of experiments where a neural network is trained to predict two different product classification outputs. An basic active learning experiment is conducted by labelling a small set of uncertainty-sampled products.

We corroborate the common finding that more input features tend to lead to better test set performance. Deep models with random embedding inputs are the most performant, which gain additional boost when augmented with embeddings from pre-trained deep models. Linear models do best with sparse 1-hot inputs, and need a wider selection of input features to do well. Multi-objective learning produces higher accuracy for objectives with few labels, but causes the training to destabilise in many cases; this is probably caused by noise in the training data and a poor choice of optimiser. Preliminary results from uncertainty sampling suggest it is capable of finding hard-to-classify products, but tends to sample products of a similar nature, and might require very small iteration sizes to be effective.

Keywords: machine learning, deep learning, transfer learning, active learning, multi-objective learning, label-efficient learning, data engineering

Acknowledgment

I would like to wholeheartedly thank my thesis examiner Magnus Boman London,
UK, July 2, 2018

Mattias Arro

Contents

List of Figures	1
1 Introduction	3
1.1 Problem	4
1.2 Purpose	4
1.3 Goals	4
1.4 Hypotheses	5
1.5 Ethical Consideration	6
1.6 Sustainability	7
1.7 Delimitations	7
1.8 Outline	7
2 Background	9
2.1 Input Representations	9
2.1.1 Categorical Input	9
2.1.2 Text Input	10
2.1.3 Image Input	11
2.2 Models Considered	11
2.2.1 General Models	11
2.2.2 Text Models	14
2.2.3 Image Models	14
2.2.4 Downstream Models	15
2.3 Combining Models	16
2.3.1 Ensembling	17
2.3.2 Joint Training	17
2.3.3 Multi-Objective Learning	18
2.4 Active Learning	18
2.5 General Practices in Machine Learning	20
2.5.1 Hyperparameter Tuning Using Bayesian Optimisation	21
3 Data Life Cycle	23
3.1 Data Preprocessing	23
3.2 System Architecture	25

CONTENTS

3.2.1	Airflow Pipelines & Pipeline Runs	26
3.2.2	Dataflow Pipelines	26
3.2.3	TensorFlow and ML Engine	30
4	Method	33
4.1	Visual Similarity	33
4.1.1	Evaluation & Results	33
4.2	Independent Models	34
4.2.1	Baseline Model: Wide & Deep	37
4.2.2	Input Representation & Model Type Combinations	38
4.2.3	Evaluation	38
4.2.4	Results	40
4.3	Multi-Objective Training	40
4.3.1	Category Structure	40
4.3.2	Labelling of Exclusive Products	44
4.3.3	Evaluation	44
4.3.4	Experiments & Results	45
4.4	Hyperparameter Selection & Tuning	49
4.5	Active Learning	50
4.5.1	Results	51
5	Discussion	53
5.1	Visual Similarity	53
5.2	Individual Models & Hyperparameters	54
5.3	Multi-Objective Training	55
5.4	Active Learning	56
6	Conclusion	57
6.1	Future Work	58
6.2	Practical Take-Aways	58
References		61
Declaration		69
Appendices		69
A	Additional Background	71
A.1	Neural Models	71
A.1.1	1D Convolutional Neural Networks	71
A.1.2	Recurrent Neural Networks	72
A.2	Unsupervised and Semi-supervised Learning	72
A.2.1	Autoencoders & Variational Autoencoders	73
A.2.2	Conditional Variational Autoencoders	74
A.2.3	Generative Adversarial Networks	74

List of Figures

2.1	Ensebling / Stacked Generalisation vs Joint Training. Boxes represent inputs (e.g. embedding vectors), hidden layers, and outputs. In ensembles and stacked generalisation (left), we have separate models with no parameter sharing, and the final output is computed either through voting or passing the output of the ensemble components through the meta learner model. In joint training (right), the two distinct inputs are concatenated and the parameters of the hidden layers are shared.	18
3.1	Dataprep: the histograms of field values of a subset of fields.	23
3.2	High-level system architecture of the ML pipeline	27
3.3	Airflow DAG: preprocess, train model, predict, update the database. Tasks with green outline have completed, others have not been scheduled yet..	27
3.4	A running Dataflow pipeline of data preprocessing	29
4.1	Nearest neighbours based on ElasticSearch title matching (TF-IDF-like)	35
4.2	Nearest neighbours based on ANN search of image embeddings	35
4.3	Nearest neighbours based on ElasticSearch title matching (TF-IDF-like)	36
4.4	Nearest neighbours based on ANN search of image embeddings	36
4.5	Nearest neighbours based on fulltext search on discretised/tokenised image embeddings as in [55]. Here scores from the title and image embedding fields was combined.	37
4.6	Input types. Orange: embedding representation, blue: 1-hot or k-hot, red: embedding extracted with a pretrained deep network.	39
4.7	Evaluation metrics of the rule-based training objective on the final time step. Minus denotes a train run that was not scheduled (inadvertently).	41
4.8	x axis = train step (batch number), y axis = training loss of rule-based objective.	41
4.9	PR AUC of all combinations of input and model types; x axis = train step (batch number), y axis = PR AUC. Refer to figure ... for exact scores.	42
4.10	Recall of all combinations of input and model types; x axis = train step (batch number), y axis = recall..	42
4.11	PR AUC on final timestep (300k, or 220k for inceptionv3)	42

4.12 Exclusive (green) vs independent (yellow) categories. A product belonging to a mutually exclusive category would not belong to any other category (e.g. a jacket is never a wine), while a product can belong to many independent categories (e.g. the same jacket can be a bomber jacket as well as a leather jacket).	43
4.13 A multi-objective model. Left: gradients (pink) when a rule-based label is present. Right: gradients (pink) when an exclusive label is present.	44
4.14 Unregularised training loss of rule-based objective. Red: only rule-based objective. Blue: until step 35K, only rule-based objective, then a combination of rule-based and exclusive. Magenta: rule-based and exclusive objectives loaded from the same files. Shuffle buffer size: 15K.	47
4.15 Unregularised training loss of rule-based objective. Blue: only rule-based objective. Red: combination of rule-based and exclusive objective. Magenta: rule-based and exclusive objectives loaded from the same files. Shuffle buffer size: 15K.	47
4.16 Validation accuracy per time step. Green: multi-objective model (exclusive, rule-based) with interleaved input files. Blue: only exclusive objective. Orange: multi-objective model (exclusive, rule-based) loaded from the same training file (unstable). Shuffle queue: 15K.	48
4.17 Validation accuracy per time step. Left green: exclusive objective in multi-objective training. Left blue: exclusive objective. Right orange: exclusivised objective. Shuffle queue: 100K.	48
4.18 Number of misclassified products in validation set.	49
4.19 Worst performers in validation set.	49
4.20 Hyperparameter ranges and the values found during automatic tuning. A single value in the “full range” column means the value was fixed.	50
4.21 Active learning: validation accuracy of the exclusive objective when training on a small number of randomly sampled vs uncertainty-sampled products.	51

Abbreviations & Definitions

AE Autoencoder

AF Airflow

AutoML a GCP product for automatically training and tuning models

BonG Bag of n-grams

BoW Bag of Words

CIFAR10 a computer vision challenge

CNN Convolutional Neural Network

CT-GAN a variant of WGAN

embedding a multi-dimensional dense vector output by a neural network or learned through some other SGD process

ES ElasticSearch

exclusive labels and training objective created by people applying manual labels (1 category per product)

exclusivised rule-based labels that have been converted to exclusive by picking a random label of an exclusive category

GAN Generative Adversarial Networks

GCP Google Cloud Platform

GCS Google Cloud Storage

GRU Gated Recurrent Units

ImageNet a large computer vision dataset and challenge

JS divergence Jensen Shannon divergence

KL divergence Kullback Leibler divergence

List of Figures

LSTM Long Short Term Memory

MF Matrix Factorisation

ML Machine Learning

MLE GCP Machine Learning Engine

MNIST a computer vision challenge consisting of hand-written digits

NLP Natural Language Processing

NN Neural Network

PCA Principal Component Analysis

RNN Recurrent Neural Network

rule-based labels and training objective produced by a deterministic keyword-matching system (1 ... n categories per product)

SGD Stochastic Gradient Descent

TF TensorFlow

TF-IDF term frequency - inverse document frequency

UUID Universally Unique Identifier

VAE Variational Autoencoder

WGAN Wasserstein Generative Adversarial Networks

1. Introduction

Machine learning (ML) has recently become a popular research area, and is increasingly applied in industry. A lot of this newfound interest and hype is directed at neural networks and deep learning. This focus is not unfounded - approaches based on deep networks continue to break benchmarks in core machine learning research areas such as computer vision [11, 53, 3], speech recognition and synthesis [23, 65], many natural language processing tasks [66, 72, 19], game play [59, 48], and even enabling novel applications such as style transfer [18] and content generation [21].

Artificial neural networks consist of layers of transformations that map multi-dimensional inputs to (usually multidimensional or structured) outputs. A single-layer neural network is a linear transformation of the inputs; each additional layer with a non-linear activation function enables the network to partition the output space and hence approximate more complex functions. Complex models are prone to overfitting and require more training data and heavier regularisation, yet deep models are somewhat unique in that their performance continues to increase with the size of the dataset, while the benefits of more data taper off for other models.

The naive conclusion is to use deep learning only when labeled data is abundant, but the depth of a model is not a binary “deep vs shallow” decision - one can start out with a shallow model and increase model complexity to see the effects on performance and generalisation. Logistic regression might well be the most appropriate model for a classification problem, but this is often not obvious up front, so it is beneficial to build models in a framework that also supports deep models. In addition to easy experimentation with model architectures, it is easy to do transfer learning with deep models on popular platforms: image or text models pretrained on large datasets can be utilised to increase product classification performance, and representations learned for classification could in turn be used by a downstream product recommender system. Datasets with abundant unlabeled data can benefit from unsupervised and semi-supervised learning using deep generative models.

Each project has different requirements on predictive performance, label complexity, interpretability, computational resources and engineering challenges. Neural networks are a good fit for this project due to their amenability to transfer learning and flexibility in handling multi-modal inputs and multi-task objectives. As neural networks are notoriously data-hungry, the question immediately becomes - how to do it without providing many labels.

1.1 Problem

The client company gathers data from various affiliate networks (that in turn get their data from various retailers) and displays it on their website. There are millions of products belonging to roughly 1300 categories, and categories follow the usual nested tree structure. The incoming data is noisy and inconsistent: what kind of data is stored in which column varies across affiliate networks, across retailers within an affiliate network, and the data within a retailer can have lots of missing values, noisy text, missing images, etc. There is currently a rule-based system for assigning products to categories. All products matching a condition (e.g. title contains the word “trousers”) will be assigned to that category, i.e. labels provided by the rule-based system are not mutually exclusive. This way of categorising products works relatively well on some categories, but such a rule-based system has several drawbacks: these rules are cumbersome to define, their evaluation is manual, they fail to match a large fraction of products that in principle should be in a given category, it is hard to trace back the rule that caused a false positive, and such rules are limited to textual data.

The client needs a classification system to replace the old way of categorising products. The predicted outputs can either be a set of independent binary classifications, or the category structure has to be re-organised to ensure classes are mutually exclusive. The system should be able to learn from the output of the old system, and if possible produce models that can be used in downstream tasks such as recommend systems and product similarity models. The highest priority is low label complexity, beating requirements for high accuracy and interpretability. The system should be robust to noisy inputs; data preprocessing should not consider the idiosyncrasies of each affiliate network. There is an additional feature the client requires: given a product image, the visitor should be shown products that are visually similar.

1.2 Purpose

The academic purpose of this work is to (1) assess the relative strengths of different kinds of models and their combinations, and (2) to determine whether an active labelling strategy reduces label complexity on a real-world dataset. Analogously, the commercial purpose is to (1) obtain a model with powerful predictive capabilities, and to (2) reduce costs by using an efficient labelling strategy, and (3) obtain a high-quality product similarity score.

1.3 Goals

The goals of the work, in chronological order, is to:

- Use a pre-trained 2-dimensional convolutional neural network (2D CNN) to extract features for an approximate nearest-neighbour search of visually sim-

1.4. HYPOTHESES

ilar products.

- Build an interface for subjectively evaluating the visual similarity algorithm.
- Train baseline model to reproduce the behaviour of the rule-based system.
- Train and evaluate a number of different models on the **rule-based labels**.
- Define a subset of the 1300 independent categories to be exclusive (non-overlapping). Build an interface for assigning **exclusive labels** to products either through **uncertainty** or **random** sampling.¹
- Train models on both the exclusive labels, and in a multi-objective manner where both the exclusive and rule-based labels are used. Evaluate whether multi-objective training improves accuracy over training just on the exclusive labels.
- Implement the active labelling mechanism defined in 4.5 and train a well-performing model in two settings: use labels obtained either through random or uncertainty sampling, and evaluate whether uncertainty sampling improves model performance.
- Document the results as well as the technical architecture and workflow.

The last four steps were done in a somewhat iterative manner. When it was evident during active labelling that a new category was needed, it was added. Training and evaluating multiple models is inherently a cyclical process, as is finding out which multi-objective learning approaches work best.

1.4 Hypotheses

The following hypotheses were postulated prior to running most experiments²:

1. Pre-trained 2D CNNs perform reasonably for visual item similarity, but fine-tuning might be needed.
2. Deep models with histogram inputs do not improve substantially on linear models.
3. Shallow models with embedding inputs generalise better, and deep models with embedding inputs generalise slightly more.
4. Pre-trained 2D CNN are consistently good predictors of clothing products, yet fail to accurately predict categories such as technology.
5. Active learning substantially decreases label complexity, however uncertainty sampling might give poor results on the first round.

How these hypotheses were evaluated is described in the relevant subsections of 4.

¹The final version of this interface was built by an employee of the client company, though earlier versions for displaying the classification / similarity results were implemented by the author.

²the Wide&Deep baseline model had been trained on rule-based labels

1.5 Ethical Consideration

Given how pervasive machine learning is becoming, it is crucial that there is discussion about the safety, transparency and bias of machine learning systems. There have already been cases of black box algorithms being used to make decisions that severely influence a person's life - for example predicting reoffending probability of convicts to determine whether they are given parole - with terrible results. The model that was touted to be an objective substitute to a subjective judge has simply learnt the biases in the data (the prejudices and racism of the judges at the time). It consistently underestimated reoffending rate of white convicts and overestimated the reoffending rate for black convicts [35]. Even though the data did not contain explicit information about race, the algorithm was capable of implicitly detecting race through some other variable that was correlated with race.

Most data scientists will never build models with such severe consequences, but even more subtle decisions made autonomously by algorithms can prolong the inequalities of our society by consistently favouring certain characteristics such as race, gender, place of origin, etc. Prime examples are algorithms that determine whether one gets a mortgage or not, what the premium on their insurance would be, whose CV gets shortlisted for a position, and so on.

The worst case scenario for inaccurate, opaque or biased product classification is embarrassment for the client company, therefore ethical considerations are not of much concern in our case. It is still worth reminding the reader that machine learning models always contain some kind of bias: from the data it uses as input (the way it is gathered, what is gathered), the way the data is processed, and the kinds of models used. One way to quantify the fairness of a binary classifier is the p% rule, which sets bounds on the ratio between the probability of the outcome given the sensitive attribute [75]. One way to overcome a classifier that seems to rely on some "nuisance parameter" (such as gender or race - even when the parameter is removed from the training data) is adversarial training [75] to ensure the model's output distribution does not depend on the sensitive parameter.

A related concern to fairness is privacy: ensuring that a ML model does not leak its training data to an attacker that might orchestrate a large series of queries on the model. Each individual query might not reveal much, but with a large number of queries the attacker can see how the statistics of the output distribution changes; paired with an external dataset for cross-referencing or knowledge that a data point was added to the training set, it is possible to recover some of the original data, though not necessarily with high fidelity or probability. A countermeasure is differential privacy, which can be applied to the data sharing process or at the algorithm level. Algorithms that are differentially private inject noise at the input, model or output level to ensure that the output distribution of the model does not change substantially with the inclusion or exclusion of a data point.

1.6 SUSTAINABILITY

1.6 Sustainability

The reality of any e-commerce company is that increase revenue almost by definition means more waste and increased carbon released into the atmosphere as a result of production and transport. There are efficiency gains (fewer bricks-and-mortar stores, fewer consumers driving to a shopping mall, increased automation), however it would be wishful thinking to assume the net effect of e-commerce on the environment is positive. Efforts to mitigate these unwelcome side effects can be successful at a government level, though the author firmly believes there ought to be stronger intergovernmental regulation to tax carbon emissions and the consumption of materials, as currently there is absolutely no incentive for retailers or consumers to reduce waste, and few incentives to recycle it.

1.7 Delimitations

As the intention was not to produce the model with the highest accuracy or absolute lowest label complexity, the selection of models and active learning scenarios was limited. Many directions for data-efficient learning are described in the theory section, but only a pragmatic subset of these is used.

1.8 Outline

The following report is organised as follows. The required theoretical background is covered in section 2: the different ways to represent input, various neural network models that could be used, joint training, ensembling and active learning. In section 3 we describe the system architecture and technical set-up, and in section 4 we describe each individual experiment separately, along with evaluation approaches and results. A more speculative and subjective analysis of the results is presented in section 5. A summary of the work is presented in section 6, along with implementation-specific tips and proposals for future work.

2. Background

In this chapter we look at the ways one can represent input data (section 2.1), and which models might be used for different input modalities (section 2.2). We consider ways to increase label efficiency with active learning (section 2.4), and general approaches to effective machine learning (section 2.5) as well as combining different models either through joint training or ensembling (section 2.3). Label efficiency can also be achieved through unsupervised or semi-supervised learning, which is described in appendix A.2 but is not used in our experiments.

2.1 Input Representations

2.1.1 Categorical Input

The simplest option for representing categorical input is 1-hot encoding, where input has as many dimensions as there are distinct categorical values (vocabulary size); a single dimension is set to 1, with all other dimensions set to 0. This is a straightforward representation: for a simple model such as logistic regression, we can clearly interpret the model parameters and see how the presence or absence of a given category increases or decreases the likelihood of a given output. However, in cases where vocabulary sizes get large, and the number of outputs (the number of units in the next layer in a deep network, or the number of output units in a shallow one) increases, this kind of encoding can really blow up the number of parameters of the model.

An alternative is to use random embeddings: dense, low-dimensional representations of a high-dimensional vectors. It has been shown in compressed sensing literature that if a high-dimensional signal in effect lies on a low-dimensional manifold, then the original signal can be reconstructed from a small number of linear measurements [7]. This has a useful implication for machine learning: categorical variables with large vocabularies of size d can be represented with random embedding vectors of size $M = \log(d)$. This is because it is possible to reconstruct any d -dimensional k -sparse signal using at most $k \log_k^d$ dimensional vectors [2], and 1-hot vectors are 1-sparse, as only one dimension is nonzero.

An intuitive example is given in [17] about natural images: a 1-million-pixel

image could have roughly 20,000 edges¹, i.e. it lies in a roughly 20,000-dimensional manifold from which the original image can be constructed with high fidelity. The 1-million-dimensional image could be projected into a M-dimensional space “by taking M measurements of the image, where each measurement consists of a weighted sum of all the pixel intensities, and allowing the weights themselves to be chosen randomly (for example, drawn independently from a Gaussian distribution)”[17]. The projection is random because the weights for the sum of pixel intensities are chosen randomly. Why the projections need to be random is motivated by another intuitive example. When light is shined on a 3-dimensional wireframe, its shadow is a 2-dimensional projection of it. The projection could lose some important information about the regional object if the direction of the light source is chosen poorly, however random directions are likely to result in projections where every link of the wire will have a corresponding nonzero length of shadow.

2.1.2 Text Input

The simplest way to represent text is bag words (**BoW**), which is simply the histogram of word occurrences in the text. BoW representations give equal weight to each of the words in the text, and the model has to learn the relative importance of each of these. A commonly used representation in information retrieval (IR) is **TF-IDF**, which stands for “term frequency - inverse document frequency”. TF-IDF multiplies the term frequency (number of times the word occurred in the text) with its inverse document frequency (the inverse of how often a occurs in all documents). This has the effect of giving lower scores for common words, and higher scores for words which appear often in a document but not too frequently in the whole corpus. Another common representation is bag of n-grams (**BonG**), which is a histogram of character n-grams.

The above representations are sparse, like 1-hot encoding of categorical variables. Text can also be represented as a sequence of **word embeddings**. Word embeddings can either be random vectors or representations learned with word co-occurrence algorithms such as word2vec [47] and GloVe [51]; these representations can remain fixed throughout the learning procedure, or updated as part of the stochastic gradient descent (SGD) when applicable. A simple way to represent text is to average the word embedding contained in it. The average could also be weighted by the TF-IDF score of each word². In [45], paragraph vectors are created in a similar manner to word2vec vectors by relying on these vectors to predict the next word in a sentence.

More complex methods use recurrent neural networks (**RNNs**) to combine a sequence of word embeddings into a fixed length vector (see section 2.2.2). These kinds of models are good at disambiguating the potentially many meanings a word might have. Exactly what gets persisted in the final vector representation of text

¹Technically: wavelet coefficients with a significant power, which roughly corresponds to a superimposition of edges.

²The author is unaware whether this has been tried before, but it seems a promising approach.

2.2. MODELS CONSIDERED

depends on how the model is trained - two RNNs trained on different tasks (such as sentiment analysis and named entity recognition) would need to store different information in its hidden state to successfully perform their relevant tasks, hence the encoding for the same sentence would be different for either model. Still, an RNN that is trained on one task could provide useful features for another. Encoding text as fixed length vectors using RNNs has been interpreted as compressed sensing [1], with such vectors being “provably at least as powerful on classification tasks, up to small error, as a linear classifier over BoNG vectors”.

2.1.3 Image Input

Images can be represented as dense 3-dimensional tensor. A 28x28 RGB image could be represented as a tensor with shape [28, 28, 3], with the final dimension corresponding to the green, red, and blue intensity values at a given x/y coordinate. These intensities are usually in the range 0 ... 255, but are normalised before input to machine learning model. Similarly to RNNs encoding a sequence of words to a vector, 2D convolutional neural networks (2D CNNs) can be used to extract dense feature vectors from raw images (image embedding), further discussed in section 2.2.3.

2.2 Models Considered

The following sections describes a selection of models that could be used for our classification task. This is by no means an exhaustive list, nor is it a selection that is expected to have the highest predictive performance individually. These models might become part of an ensemble or be trained jointly with other models, where diversity matters much more than the performance of any individual model. Engineering considerations and time constraints also influence the selection of models. We did not want to spend time implementing nontrivial models from scratch, or to use many frameworks for training models. TensorFlow is a widely used machine learning library with a good selection of open source models, in particular a selection of pre-trained models for computer vision and NLP, and has arguably the most mature deployment ecosystem. Therefore models that did not have an open source TensorFlow implementation or a trivial implementation were excluded.

Section 2.2.1 describes models that can take categorical and text inputs, section 2.2.2 looks at some neural models that take only text inputs, and section 2.2.3 describes 2D CNNs that can classify or extract features from images.

2.2.1 General Models

Logistic Regression

Logistic regression is a linear, discriminative classifier that is a common baseline model, since it is both interpretable and efficient to train. There are methods that

take time linear in the number of non-zeros in the dataset, which is the smallest amount possible, and it can be made to handle non-linear decision boundaries by using kernels [38].

We have two cases: for independent categorical outputs we use a separate binomial logistic regression model for each output class, and for mutually exclusive classes we use multinomial logistic regression. This is shown formally in eq 2.1 and 2.2, where where X and W correspond to the input and weight matrices, respectively.

We follow the notational trick where the first row of X is always 1, and the first row of W corresponds to the bias term. The loss function of this model as well as how it is optimised is described in section 2.2.1.

$$p(y|X, W) = \text{sigmoid}(W^T X) \quad (2.1)$$

$$p(y|X, W) = \text{softmax}(W^T X) \quad (2.2)$$

Feedforward Neural Network

Feedforward neural networks (also called deep feedforward networks, multi-layer perceptrons) have been described as function approximator, whose goal is to approximate some functions f^* . We expect some familiarity with neural networks from the reader; for an excellent overview of the various use cases and models of deep learning, refer to [20]. In our case, the input is some representation of categorical, text or image input - often an embedding extracted with another model or assigned randomly. The hidden layers of our network are homogenous: they use the same activation function (ReLU, sigmoid, or tanh) and have the same number of units; activation functions and the number of units in hidden layers are determined through hyperparameter tuning. Similarly to logistic regression, we have two cases: sigmoid activation for independent categories and softmax layer for mutually exclusive categories.

Wide & Deep

Wide and deep model was originally proposed for recommender systems [9], but can just as well be used for classification. In this model a deep network and a linear logistic regression model are trained jointly in the same SGD learning process, rather than trained separately and then ensembled. The wide component is good at remembering “feature interactions through a wide set of cross-product feature transformation” while the deep model provides good generalisation. In our experiments we did not use cross products of features.

Loss Function and Optimizers

Logistic regression and neural networks are optimised with some variant of gradient descent. The loss function was the same for multi-label and multi-class training

2.2. MODELS CONSIDERED

objectives: binary cross-entropy between the training data and the model distribution, with the loss value averaged across all classes. This is given in eq. 2.3, where θ corresponds to the model parameters such as the weights and biases of the model, C corresponds to the number of classes, and $P(c = 1|x, \theta)$ gives the predicted probability that the item belongs to class c given the feature vector x .

$$NLL(\theta) = \sum_{c=1}^C \sum_{i=1}^N [c_i \log P(c = 1|x, \theta) + (1 - c_i) \log(P(c = 0|x, \theta))] \quad (2.3)$$

We are minimising the negative log likelihood (NLL) rather than maximising the product of likelihoods. Multiplying large numbers of probabilities could result in numerical underflow and rounding errors, therefore it is pragmatic to work with sums of log probabilities instead. The Adam [40] optimiser works very well with default hyperparameters, and is therefore used for all stochastic gradient descent updates.

Cross-entropy is a standard loss function for multi-class and multi-label classification, however in our cases it may have downsides. There is at times ambiguity as to which category a product should belong, therefore even hand-assigned labels are somewhat arbitrary. The similarity between categories is ignored by this loss function: a short coat that is misclassified as an earring has the same loss value as a coat that is classified as a jacket. Additionally, the rule-based labels are incomplete: the lack of a label does not always mean that the product does not belong to the given category, yet the model gets penalised for cases where such label is missing but the model correctly predicted a semantically similar class.

The Wasserstein distance metric (also called earth mover distance) can be used instead, which gives the cost of the optimal transport plan for moving the mass of one probability distribution to another. In our case the two distributions are the label distribution of a data point (1-hot / k-hot vector) and the class distribution predicted by the model. The optimal transport plan is weighted by a $K \times K$ distance matrix, where K is the number of classes and the entries in the matrix correspond to how dissimilar the classes are³. Intuitively, transporting probability mass from the category “Coats” to “Earrings” should be more expensive than from “Coats” to “Jackets”. As there could be many acceptable transport plans that preserve the probability mass, an iterative algorithm needs to be used to find an approximately optimal one that also considers the different cost associated with transporting probability from one class to another. In [16] the Sinkhorn-Knopp algorithm with an entropic regularisation term is used as a relaxation of the exact but costly linear programming solution. This is still an expensive operation compared to cross-entropy or mean squared error, but should be manageable in the 1000-dimensional space of categories, as opposed to 500x500-dimensional image space where it is often also applied.

³Similarity is known *a priori*, e.g. derived from WordNet hierarchy. It would be interesting to explore similarity matrices created using cosine distance of embedding vectors of categories; embeddings could be extracted with an RNN from the description of the category

2.2.2 Text Models

1D Convolutional Neural Networks

1D CNNs are a simple yet powerful family of neural models that can work on variable-length sequences of data. These models convolve several “filters” on top of the original sequence, which produce variable-length activations; some (usually max) pooling operation ensures the final activation is of a fixed length. Therefore such models can be used as classifiers as well as feature extractors. We do not use 1D CNNs in our experiments; refer to appendix A.1.1 for a more thorough overview.

Recurrent Neural Networks

Recurrent neural networks (RNNs) are another family of models capable of processing sequential data. RNNs take a sequence of inputs (e.g. dense word embeddings), and either output a prediction after each time step, or after processing the whole sequence. RNNs and its variants maintain a memory vector that represents the input the model has seen so far; the state of the hidden vector after processing all data points is an embedding of the sequence, which can be used for prediction or in downstream models. We do not use RNNs in our experiments; refer to appendix A.1.2 for a more thorough overview.

Deep Averaging Network

Deep averaging networks are simple architectures where the embeddings of the input tokens (words, in our case) are averaged and fed through a deep network [33]. In fact, a deep model with embedding inputs is very similar to a deep averaging networks, the main difference is in the way the model is optimised and regularised. We use the Universal Sentence Encoder [8] in our experiments, which is a deep averaging network pre-trained on a large text corpus.

2.2.3 Image Models

2D CNNs are important historically - excellent results in computer vision revived interest and funding in deep learning - but also applicable to a wide range of problems. Computer vision CNNs tend to be complex models with millions of parameters and hundreds of layers. It is very compute-intensive and time-consuming to search for new architectures, and the analysis of these is beyond the scope of this work. Our goal is to use a CNN that has a reasonable performance with decent speed.

The general architecture of a 2D CNN is as follows. Patches of feature extractors are convolved over the x and y axis of the image; these are often small, e.g. $3 \times 3 \times c$ patches that span all the c channels of the image ($c = 3$ for RGB images), and convolution could be strided. Former models used larger patch sizes, but stacking several smaller patches could achieve similar receptive fields while allowing for a larger number of non-linearities for the same number of parameters. One or more layers

2.2. MODELS CONSIDERED

of convolutions would be followed by a (often max) pooling layer. Computer vision architectures tend to re-use the same sub-structure repeatedly: either the same kinds of convolutions followed by pooling, or the “network in a network” structure popularised by GoogLeNet [62]. After a number of such repeated structures, there are normally a few fully-connected layers followed by a softmax or sigmoid layer.

Transfer learning is particularly applicable to computer vision. Successful computer vision models require a notoriously large training datasets, which would be hard to come by for every problem. Instead, models are often initialised to the values of a model that has been trained on a large dataset such as ImageNet [54], excluding only the last layer(s) of the model that has learned features that are specific to the original problem. Representations learned at the lower layers are remarkably universal and useful for other tasks. Therefore pre-trained 2D CNNs can be used as feature extractors for various other tasks, or the models could be fine-tuned to learn representations that are even more useful to the new task.

2.2.4 Downstream Models

Item Similarity

Similarity measures based on tokenised text are easily accessible to e-commerce companies. The full-text search engine Lucene is at the core of the popular open source NoSQL database ElasticSearch (ES). ES provides document similarity metrics that are some combination of normalised TF-IDF vectors between all documents and a query, which could be another document. This gives decent results in many cases, but it does not take advantage of a product image.

Embeddings extracted with a deep network can be viewed as a distributed representation of the original data, which carries semantic meaning [27]. In the contrived examples of distributed representations, each dimension would represent the degree to which a property (e.g. of shape or colour) is present of an object, however the dimensions are not necessarily so nicely disentangled in the representations learned by deep networks⁴. Still, semantically similar data points will have similar embedding vectors, which help shallow models make accurate predictions.

Therefore we can compare any two embedding vectors extracted with the same neural network using standard vector space similarity measures, such as cosine similarity ($1 - \text{normalised dot product}$ between the vectors). Such embeddings can be extracted from product images with 2D CNNs pre-trained on another task, 1D CNNs or RNNs on product title and descriptions, or from the joint embeddings described in section 2.3. With large numbers of products, calculating pairwise similarity scores becomes intractable, therefore approximate nearest neighbour methods could be used [4].

⁴The author has not seen strong evidence in literature to support either case, but at least in [30] extra steps had to be taken to *prevent* certain latent variables from encoding properties of the data that were designed to be captured by other latent variables. This implies that by default, each latent variable contains a little bit of information about each aspect of the input, and that the whole pattern of the embedding matters for conveying semantic information

Another interesting method of computing similarity scores between documents utilises the full-text search capability of tools such as Lucene, which would be valuable to firms already using such technologies. The embedding vectors of products could be tokenised, by converting each dimension of each embedding vector into tokens that represent the feature with some precision [55]. For example, we could consider two levels of precision and divide the normalised feature into 10 intervals of size 0.1 and 100 intervals of size 0.01; each product would get two tokens per embedding dimension. It would be a good idea to also have some wider overlapping intervals as well, e.g. 0 ... 0.5. For example, if the first embedding dimension is 0.46, the inserted tokens would be d1_0.4 and d1_0.46. Full-text search engines would assign higher scores to co-occurrences in the more precise tokens and lower scores to co-occurrences in more coarse token.

Recommender Systems

Several methods have been proposed that use dense embeddings learned by deep models as part of a recommender algorithm. These are briefly described here to motivate our focus on models that obtain dense embeddings of products.

Models that deal exclusively with deep embeddings [14] require huge amounts of user-product feedback pairs, therefore the more interesting solutions are ones which work together with classical methods such as matrix factorisation (MF). MF finds - purely from the implicit or explicit feedback users give to items - vectors of latent factors that represent each user's preferences and each item's "qualities"; the inner product between a user's and item's latent factors gives a score of compatibility between the two. Deep models can augment the item latent factors by injecting the embeddings learned for another task. The main difference in these methods is how the deep embeddings are merged into MF, and how the deep embeddings are obtained in the first place. In [67], stacked denoising autoencoders are used for unsupervised feature learning of item data, while [25] uses simply visual embeddings extracted from a 2D CNN pre-trained on ImageNet.

2.3 Combining Models

Individual models can be combined to achieve greater accuracy and lower label complexity; the models can either become a part of an ensemble (described in section 2.3.1), or trained jointly as part of the same gradient descent optimisation process (described in section 2.3.2). A neural network can be optimised simultaneously for many objectives; additional training objectives can improve the accuracy of other objectives. Multi-objective learning (described in section 4.3) can be seen as a way of combining the models of individual training objectives.

2.3. COMBINING MODELS

2.3.1 Ensembling

Some ensembling approaches are expected to consist of **weak learners** or models from the same model family. Bootstrap aggregation (a.k.a bagging) draws several bootstrap samples from the training data, trains a separate model per bootstrap sample, and averages the predictions of these individual models. This reduces the variance of predictions without increasing its bias and generally leads to more accurate predictors [5]. Boosting is another common meta-algorithm that iteratively trains an ensemble of weak models, such that the data points that incurred a higher error on the previous ensemble are weighted more heavily, and each new weak model is added to the ensemble with a weight proportional to the weak learners accuracy. A common boosting algorithm is AdaBoost [15].

We are interested in ways of combining models that are separately trained and combined into a single predictor. Committee methods and stacked generalisation are more appropriate for **ensembling strong, independently trained models**. Simple solutions are voting or (weighted) averaging, where the weights can be found with k-fold cross-validation. A more appealing approach is stacked generalisation, where a meta-learner is trained on the outputs of the individual models. This meta-learner could be a relatively simple model, such as logistic regression or a decision tree, and the results of the meta-learner can be therefore quite interpretable. Refer to figure 2.1 for a schema of an ensemble of two neural networks.

2.3.2 Joint Training

All models described in section 2.2 are trained using stochastic gradient descent, so in principle it would be possible to learn the parameters of all components (1D CNN, 2D CNN, RNN, DNN, logistic regression) as part of a single optimisation loop. This would be an engineering challenge due to the different ways these model require their input to be represented and the large size of the combined model. In our case this would be clearly an overkill, but it would be an interesting research problem to examine whether joint training has advantages over ensembling in cases where we have different views of the data (in terms of considering different input dimensions or processing the inputs differently). As mentioned in the Wide & Deep paper [9], the wide component of the model has to just handle cases where the deep component falls short, and can therefore be simpler than it would need to be in an ensemble; analogously, jointly training different neural networks would allow each component to be simpler.

One of the motivations for picking mostly deep neural models was their transfer learning capability. They all produce a dense embedding vector from which separate logistic or multinomial regressions predict the class. While it is useful to have separate embeddings for images, text and categorical variables, we would also like to have a compound embedding of each product. A simple concatenation of all individual embeddings could be a useful (albeit a quite high-dimensional) representation to be used in downstream models. Alternatively, we could use this concatenation

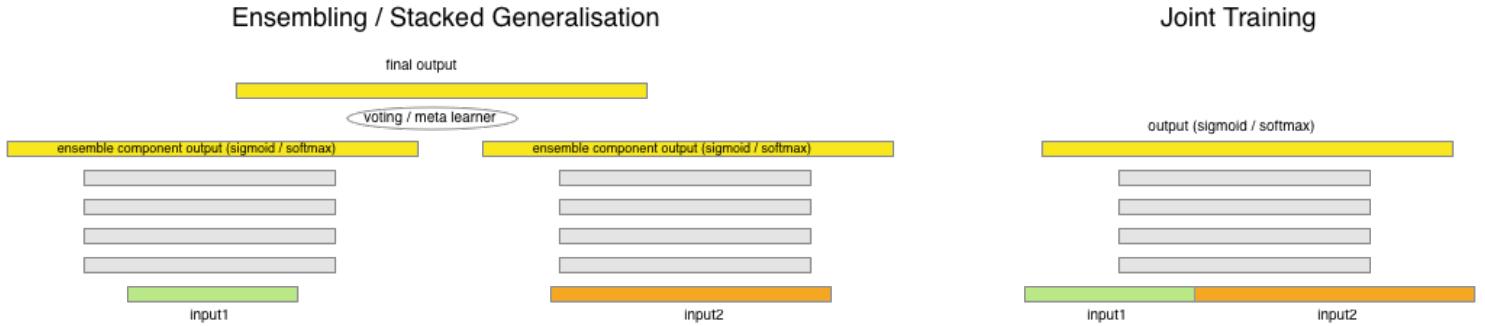


Figure 2.1: Ensembling / Stacked Generalisation vs Joint Training. Boxes represent inputs (e.g. embedding vectors), hidden layers, and outputs. In ensembles and stacked generalisation (left), we have separate models with no parameter sharing, and the final output is computed either through voting or passing the output of the ensemble components through the meta learner model. In joint training (right), the two distinct inputs are concatenated and the parameters of the hidden layers are shared.

as an input to a neural network, whose goal is to (1) reduce its dimensionality by removing redundant information and (2) predict the output classes from this compressed embedding. The latter case could act as a replacement for the meta-learner, with the difference that it takes the penultimate (as opposed to the final) layer of each individual model as input. Refer to figure 2.1 for a schema of joint training.

2.3.3 Multi-Objective Learning

A neural network can be trained to predict multiple targets from the inputs. In our setting we can consider the rule-based multi-label objective and the hand-assigned exclusive labels as separate training objectives; we might add additional training objectives that predict colour (multi-label) or conversion (binary). When using feedforward networks for multi-objective learning, the lower layers are usually shared among the objectives, and the top layers are specific to each training objective. It has been shown repeatedly that when the training objectives are somewhat related, then adding an additional objective increases the prediction accuracy across the other objectives [56, 70]. When implementing multi-objective training, care must be taken not to influence the parameters that are specific to another objective. Refer to figure 4.13 for a schema of a deep network with two objectives.

2.4 Active Learning

This section describes the main approaches to active learning; details analysis is beyond the scope of this report.

2.4. ACTIVE LEARNING

The goal of active learning is to reduce the number of labels needed to effectively train a machine learning model by being selective about which data points to label. The three general scenarios: membership query synthesis, stream-based selective sampling, and pool-based sampling [57]. In the first scenario, the algorithm synthesises a datapoint and asks a label for it. In stream-based selective sampling, an instance is sampled and the model decides (while having access to the data point's features) whether to acquire a label or not. In pool-based sampling, the model chooses which data point to label from the pool of all unlabelled samples. Our use case is a version of pool-based sampling, where the algorithm picks out a batch of products to be labelled. The rest of the section describes different approaches to picking the data points that are expected to help the model learn.

Uncertainty sampling is the most popular choice: if the model is uncertain about a prediction, obtaining a label for that data point would help it differentiate between different targets. The most common way to quantify uncertainty is by entropy of the model's predictions:

$$x_H^* = \operatorname{argmax}_x - \sum_i P_\theta(y_i|x) \log P_\theta(y_i|x), \quad (2.4)$$

where x_H^* corresponds to the most informative instance according the entropy (H) measure, and y_i corresponds to each class.

Query by committee uses a number of models trained on the labelled data which all make a prediction on the unlabelled data. Data points with highest disagreement are expected to be most informative, as by definition a larger number of models would have to be wrong in their predictions. The models in the committee do not have to be of different type like is our case, but could be e.g. a set of linear classifiers where each committee member just has different parameter values, yet each member would have to be consistent with the current set of labels (not contradict it with its predictions). Disagreement between committee members can be quantified with vote entropy [12]:

$$x_{VE}^* = \operatorname{argmax}_x - \sum_i \frac{V(y_i)}{C} \log \frac{V(y_i)}{C}, \quad (2.5)$$

where $V(y_i)$ is the number of votes given to class i and C is committee size. Alternatively, KL divergence between each committee member's predictions and the consensus P_ζ (average prediction of committee members) could be used:

$$x_{KL}^* = \operatorname{argmax}_x \frac{1}{C} \sum_i D(P_{\theta(C)} \| P_\zeta) \quad (2.6)$$

$$P_\zeta(y_i|x) = \frac{1}{C} \sum_{c=1}^C P_{\theta(c)}(y_i|x) \quad (2.7)$$

$$D(P_{\theta(C)} \| P_\zeta) = \sum_i P_{\theta(C)}(y_i|x) \log \frac{P_{\theta(C)}(y_i|x)}{P_\zeta(y_i|x)} \quad (2.8)$$

Expected model change calculates how much a model would change if we knew its label. In gradient-based learning algorithms it is possible to compute the L2 norm of the gradient vector for each combination of labelings of the unlabelled product [58], which is a direct measure of how much the model would change.

Expected error reduction considers how much the test set error is likely to decrease with the acquisition of a given label. This is done by re-training the model once per each combination of the label values and observing how either the risk or expected log loss of the unlabelled data points decreases. This is not a suitable use case for deep networks, which are trained over several epochs of the data; there is no efficient way to incrementally train the model, and when re-training with a single additional label, the change in the model is probably higher from the stochasticity in the training procedure rather than the additional label. In any case, this method is very expensive computationally.

Variance reduction tries to pick data points which are expected to decrease the variance in predictions. To do that, the Fisher information of the model parameters should be maximised. Unfortunately this involves inverting a $K \times K$ covariance matrix, where K is the number of parameters in the model. This is impractical for deep networks with millions of parameters.

Information density measures consider data points that are not just uncertain but also representative of the underlying distribution. This avoids the problem with many other approaches that ask labels for samples that are controversial, but could otherwise be outliers and not the most useful for good generalisation. In this framework, a base informativeness measure $\phi_A(x)$ such as uncertainty or disagreement is weighted by the average distance of the data point x to other data points in the distribution:

$$x_{ID}^* = \operatorname{argmax}_x \phi_A(x) \times \left(\frac{1}{U} \sum_{u=1}^U sim(x, x^{(u)}) \right)^\beta, \quad (2.9)$$

where sim is some similarity function and β controls the importance of the density term.

2.5 General Practices in Machine Learning

We briefly list some techniques and approaches that could be used for many kinds of machine learning problems. Some of these will be used in our experiments, yet some may not be pragmatic due to long train times and time constraints.

Bootstrap sampling, where data points are sampled from training data with replacement, can be used to repeatedly train the same model, and to observe the variance of its predictions. Bumping can be used to train several models (of the same family) on bootstrap samples to move around in the model space, and will pick the model that best fits the training data. This helps it explore a wider selection of models, but given that the original training data often also included as one of the

2.5. GENERAL PRACTICES IN MACHINE LEARNING

bootstrap samples, this method can still pick the model original model if it happens to have the best accuracy.

There are some versatile methods that can help regularise a neural network, or to speed up convergence by improving gradient updates. L2 regularisation should always be used to decrease moral complexity and impose prior on the model parameters, which implies that very high and very low values are unlikely. Dropout can mitigate overfitting by preventing neurons from co-adapting, encouraging each to learn representations that are useful independently [60]; this has been interpreted as training and exponential (in the number of parameters) number of models and averaging their predictions. Batch normalisation is known to stabilise training and hence speed up convergence by normalising each input to have unit variance and zero mean [32]. The effect of batch normalisation depends on the minibatch size, so layer normalisation may be used instead “by computing the mean and variance used for normalization from all of the summed inputs to the neurons in a layer on a single training case” [36]. Gradient clipping has been proposed to alleviate the vanishing and exploding gradient problem notorious in RNNs [50].

2.5.1 Hyperparameter Tuning Using Bayesian Optimisation

There are roughly 4 ways of choosing hyperparameters (parameters that are not explicitly learned from the data): manual, grid search, random search and (Bayesian) optimisation. Some manual choices are done in any case, as one still needs to decide on the intervals for the hyperparameters. Grid search, where all combinations of (discretised intervals of) hyperparameters is tried, is not scaleable, as the number of possible sets of hyperparameters grows combinatorially; a poor choice of intervals for continuous hyperparameters might miss the “sweet spot” for the given parameter, and the model would still need to be evaluated several times for a hyperparameter which has nearly no effect on model performance. In random search, sets of hyperparameters are sampled uniformly randomly from the given ranges, which is a more efficient use of compute power. Since each hyperparameter has a distinct value at each train run, this reveals more about how the model’s performance changes with respect to all hyperparameters.

The fourth option considers the set of hyperparameters as an input and the model’s validation performance as an output to be optimised. As sampling from this objective function is very expensive - it involves training a neural network until convergence - we need an approach that can explore the space of values efficiently. We do not have a closed form description of the objective function or gradients of the objective with respect to the inputs, but we can make assumptions about the objective function that can be encoded as priors in a Bayesian framework, e.g. by assuming it has a Gaussian distribution. To choose sets of hyperparameters that are likely to result in a good performance, the tuning algorithm can use the Bayes rule to compute the posterior distribution of the objective function; points that have a high posterior mean and variance should be explored [6], as these are the uncertain points that are expected to have the highest value.

3. Data Life Cycle

Section 3.1 gives an overview of how input data was pre-processed. Section 3.2 describes the technical set up required for running all these experiments and deploying the model to production.

3.1 Data Preprocessing

There were around a dozen product features that affiliate networks provided. Most of these were either categorical or textual, with just a single numerical feature. Initially, the data was analysed using Dataprep, a Google Cloud Platform (GCP) product for data wrangling, which at the time of use was in beta stage. Dataprep was used to process a sample of 800 000 products; it produced histograms of the values present in each feature column (see figure 3.1).

The histograms revealed that a lot of the input features were mostly empty, but also that many of the inputs that would naturally be considered categorical had much more unique value in them than one might expect. For example, each affiliate gives us the textual representation what they consider to be the category of the product, but rather than containing a small number of unique tokens, these



Figure 3.1: Dataprep: the histograms of field values of a subset of fields.

contained all the full category paths along with the category delimiters, which varied retailer by retailer (e.g. it was common to see both “Shoes > Sneakers” and “Shoes // Sneakers”). Representing these as categorical variables would have blown up the input space, which would have resulted in more parameters, each parameter having fewer examples to learn from. Therefore, many such “categorical” features were actually represented as text, which were tokenised and cleaned appropriately, allowing for better generalisation and smaller models.

There was a single numerical field: price. This could have been min-max normalised to the range 0 ... 1, however there was a small number of very high values that would have squashed nearly all the other prices. Rather than carefully considering how to mitigate this, the input dimension was dropped, because it is not likely to have much predictive value for product classification. It would be trivial to bring this feature back for a training objective for which it would be much more useful.

A trickier question was how many distinct tokens or categorical values to keep per input column. Keeping all of them would not have been sensible: there were still large numbers of tokens that appeared only once, often because there were some unwanted formatting characters, misspellings, or incorrect punctuation that caused a token to be considered a separate entity. The final number of tokens to keep was picked based on the histogram of tokenised values in each column; the threshold was picked roughly such that tokens that appear more than 5 times would be kept.

There was a single configuration file that dictated which models used which features as input, whether those inputs were represented as textual, categorical, or dense values; it also determined the maximum number of unique values/tokens, and the dimensionality of the embedding. This configuration file was read by Dataflow during pre-processing and by TensorFlow during inference and training, which made experimentation with different types of models and input representations considerably easier.

Below is a list of input features with information about how they were represented; it also lists the dimensionality of embeddings for the models which encoded categorical variables as embedding.

- title - text - max 8000 unique tokens
- brand - categorical - max 5000 unique values - 10 embedding dimensions
- category - categorical - max 950 unique values - 6 embedding dimensions
- rawCategory - text - max 1000 unique tokens
- description - text - max 8000 unique tokens
- gender - categorical - take all unique tokens
- size - categorical - max 100 unique tokens
- image - dense vector of 2048 or 1280 dimensions extracted with a 2D CNN

3.2. SYSTEM ARCHITECTURE

3.2 System Architecture

The following technologies were used to build the system which had to interact with existing services at the client company. Figure 3.2 shows the how data is passed between the main services, and how services and technologies interact.

- Apache Airflow (AF) - a Python framework for defining workflows of long-running tasks and dependencies between these tasks.
- TensorFlow (TF) - ML framework for Python, capable of defining many kinds of models as a computation graph, and executing this graph locally or in a distributed manner.
- ML Engine (MLE) - a GCP service for running TensorFlow models (training, hyperparameter tuning, inference).
- Apache Beam - a data processing engine akin to Apache Spark.
- Dataflow - a GCP service for executing Apache Beam workloads.
- Tensorflow Transform - a Python library with a small set of operations for data preprocessing that can run inside a TensorFlow graph as well as an Apache Beam pipeline.
- TensorBoard - a web interface for visualising the variables and metrics a TensorFlow train run produces.
- Google Cloud Storage (GCS) - Google Cloud Platform (GCP) object storage similar to Amazon S3.
- ElasticSearch (ES) - a NoSQL database with powerful full-text search and querying capabilities.
- RabbitMQ - a message queue, used for transferring data among our microservices (using the Logstash adapter, that can read from and write to (among other things) ElasticSearch and RabbitMQ).
- Flask - a simple backend web framework for Python.
- Node.js - a JavaScript backend web framework.
- React.js - a JavaScript frontend framework.
- Redux - a framework for persisting user interface (UI) state and application data in single page applications.
- GraphQL - a query language for building flexible APIs

All product data is stored in ElasticSearch (ES): the labels, top predictions of the ML system, and evaluation metrics from various train runs. ES is accessed from the public web application via GraphQL and the ML administration web UI (further referred to just as web UI¹). Data is pulled into the ML pipelines by dumping the

¹The web UI was initially built with Flask and React by the author as a quick way to get insight about the model, and then re-written as a more feature-rich version by an employee of the

results of an ES query to a local file, which is uploaded to GCS. Updates to the ES index are not done directly, since indexing the updated products is computationally expensive; instead, updates are put on a RabbitMQ queue, which is consumed by Logstash, which in turn updates products in ES at a rate that will not overburden the servers.

3.2.1 Airflow Pipelines & Pipeline Runs

All ML training and prediction happens in a batch-oriented way, encapsulated as Airflow pipelines. Each pipeline is a directed acyclic graph of tasks, where a task can be a shell command or Python function; a pipeline defines dependencies of task execution, which allows us to coordinate a series of operations that could be executed locally (inside the Docker container running AF) or remotely (such as in a GCP service). A typical pipeline (shown in figure 3.3) dumps data locally, uploads it to GCS, schedules a Dataflow job to preprocess data, polls the Dataflow service until the Dataflow job is finished, schedules an ML engine job, polls the MLE service until it has completed, and runs an update process that reads the predictions and evaluation statistics from GCS and sends the updates to the RabbitMQ queue. Reading updates and sending these to RabbitMQ is done in a parallelised manner (using multiprocessing), since the update process is bottlenecked by network latency as well as computing the appropriate category path for each product (explained in 4.3.1).

All pipelines create their own set of directories in GCS; the directory name becomes the *run_id*. Raw data is dumped into that directory; dataflow reads from and writes to subdirectories of the same run. The Tensorflow trainer also needs the *run_id* to find locations of the Tensorflow Transform model function (described below) and .tfrecords input files. A given run usually has a single model that is trained, however at times we train several models with different architectures or subsets of labels.

3.2.2 Dataflow Pipelines

There are two types of Dataflow pipelines: for preprocessing training data and for calculating product-product visual similarity. Omitting various details, dead-ends and workarounds that were needed due to technical limitations and prior system architecture choices², the pipelines had the following tasks:

Preprocessing

This pipeline loads the products dumped from ES (as JSON) and preprocesses data according to a configuration file. Figure 3.4 shows a screenshot of this pipeline, as visualised by the GCP UI.

^{client company with Node.js, GraphQL, React and Redux.}

²which were completely reasonable at the time, when the ML system was not a consideration

3.2. SYSTEM ARCHITECTURE

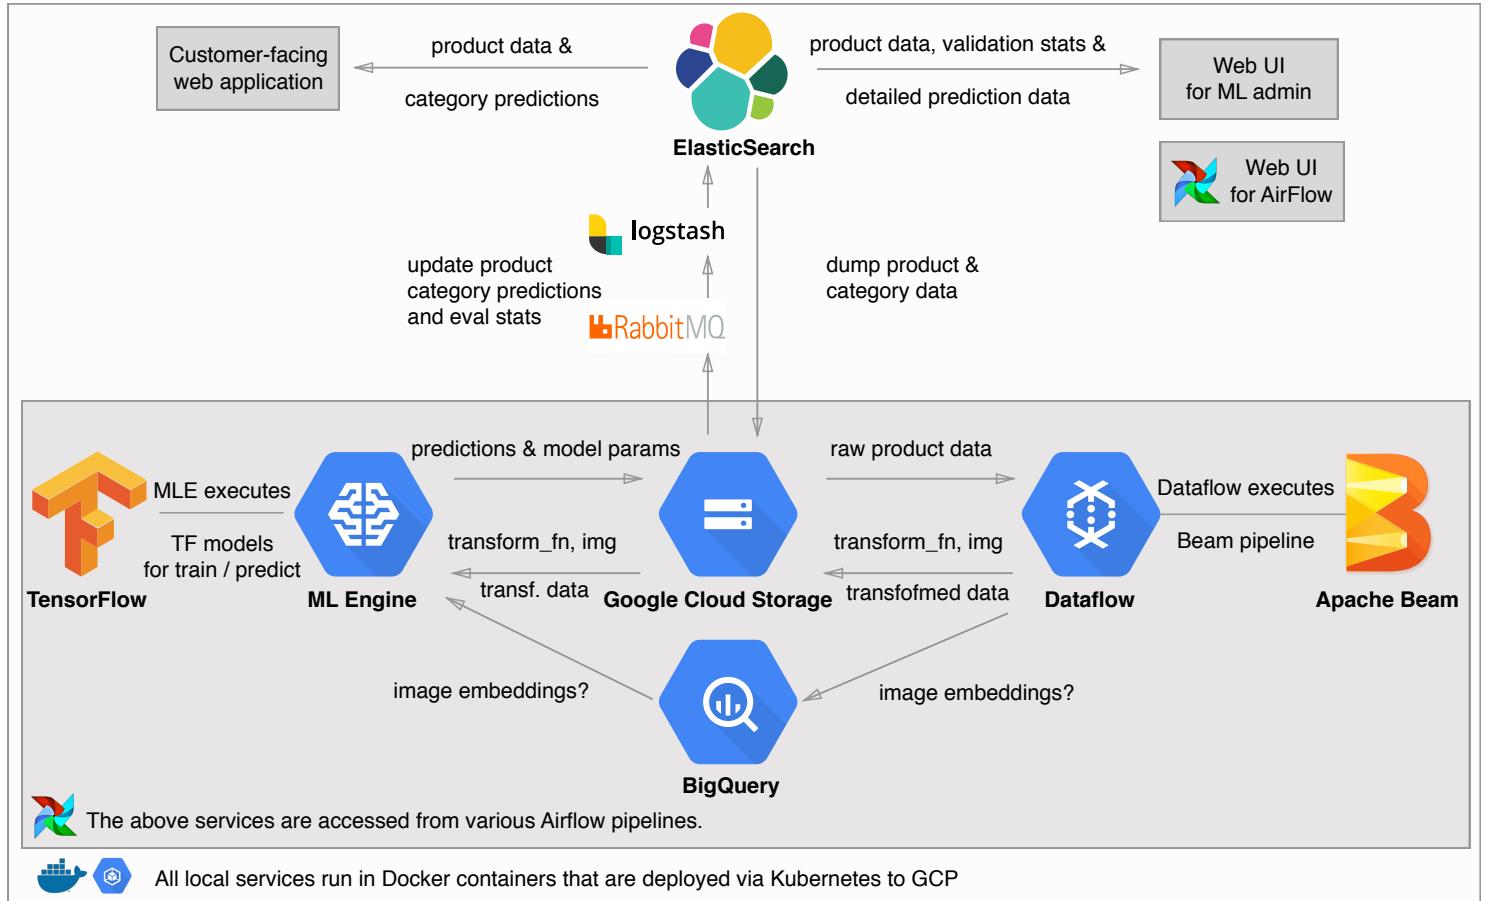


Figure 3.2: High-level system architecture of the ML pipeline



Figure 3.3: Airflow DAG: preprocess, train model, predict, update the database. Tasks with green outline have completed, others have not been scheduled yet..

In this pipeline, all fields are cleaned of obvious noise and superfluous whitespace. Text and categorical fields are tokenised and mapped to integer indices, keeping only the top k values and also computing TF-IDF scores for text fields. This is handled by TensorFlow Transform, which persists this token-to-index mapping in a **transform function** that is saved to GCS at the end of the pipeline. The transform function can be used by TensorFlow or Dataflow to convert raw text inputs to a sparse inputs; separate pre-processing run will generate a different transform function, with mostly the same tokens mapping to different indices, as the order in which they will be encountered will be different. The tokenised data is stored to GCS as a sharded .tfrecords files (with one set of files per training objective), a binary file format that TF Transform helps produce; alternatively the raw data could be saved (e.g. as JSON or inside BigQuery) and the transform function inside a TF graph could re-tokenise this raw data at runtime.

The pipeline is also responsible for downloading product images, which are saved to GCS as individual files to avoid hitting the content delivery network that stores product images multiple times. Images are a major bottleneck due to network latency. To avoid a TF / MLE job from making separate requests to fetch each file individually³, Dataflow saves the raw image content into the .tfrecords file; this increases Dataflow execution time, but reduces time and cost overall, as Dataflow is cheaper and easier to parallelise than a MLE job.

A former version of the pipeline also computed image embeddings inside the Dataflow job, but this was inefficient. We also wanted to have the flexibility to try different models for computing embeddings (e.g. Inception V3, MobileNet, AutoML models) and fine-tuning image models. Attempts to persist these pre-computed embeddings had severe overheads, and currently the approach is to re-compute these at every time inside MLE, which can leverage a GPU to do it efficiently.

Visual Similarity

The Dataflow pipeline implemented for the experiments described in 4.1 relied on image embeddings computed by the preprocessing pipeline. As explained in section 4.1, it needs to find the $k=100$ nearest neighbours of each product based on the cosine similarity of their image embeddings. The product-product similarities are the simplest case computed within products that belong to same second level category. The pipeline read the files output by the preprocessing stage, and partitioned it into datasets by the category predicted in the previous train run, as the new predictions would not be saved in the file. The per-category image embeddings were merged (using *reduce*) into matrices of image embeddings in a given category, and the top 100 approximate nearest neighbours were extracted for each product using nmslib [4] (as a *map* operation over all the per-category embedding matrices). The nearest neighbour UUIDs were persisted to GCS and a downstream AF task picked these up and updated ES with these.

³Each request has latency overhead, which compounds quickly when doing million of these.

3.2. SYSTEM ARCHITECTURE

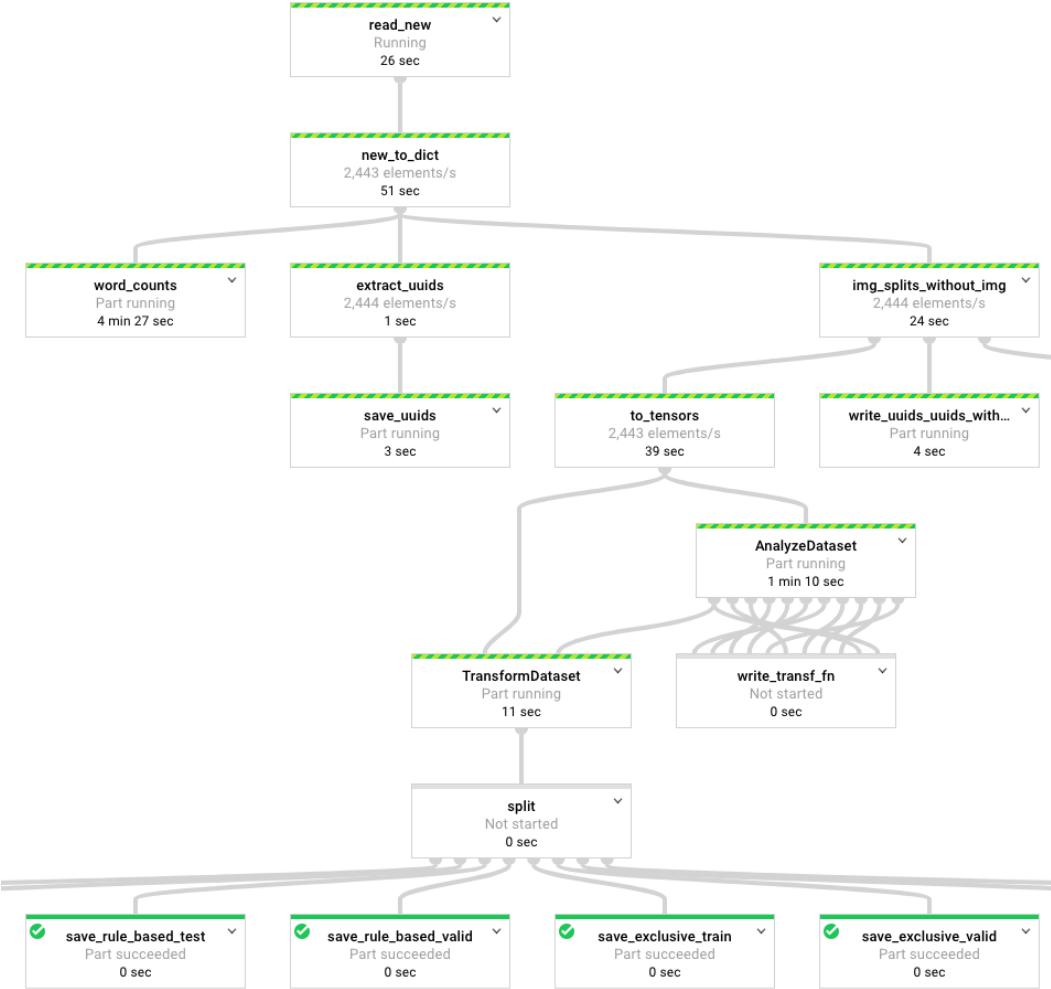


Figure 3.4: A running Dataflow pipeline of data preprocessing

This set-up will be replaced in the future. Computing embeddings in Dataflow is slow, and we may want to enable more flexible ways of restricting the subset of products that are considered (e.g. not just 2nd level category, but a 3rd level category or even globally). The most flexible approach is one in which nearest neighbours are computed in real time when a product is viewed by a user. The biggest challenge in such a system is keeping all the image embeddings in memory, which would require 32 GB of memory for the 4 million products we currently have⁴, and the number of products is likely to increase in the future.

An option worth considering is deploying this as a TensorFlow model to MLE, which reads the embedding matrix in as a sharded tf.Variable; TF can operate on

⁴ $4000000 * 2048 * 32 = \text{number of products} * \text{embedding dimensionality} * 32 \text{ bits per number}$

different shards in parallel, which reduces latency. The deployed model can store one embedding matrix per category, which would result in matrices with (hundreds of) thousands of rows. Performing a vector-matrix product at this scale is fast enough for real-time use⁵. We can also use autoscaling built into MLE that would increase the number of workers to handle high loads, which would also put all workers to sleep after 15 minutes of inactivity.

3.2.3 TensorFlow and ML Engine

All models were implemented using TensorFlow, using higher level APIs (such as `tf.data`, `tf.estimator` and `tf.learn`) when possible. This was somewhat challenging, as the high-level APIs were poorly documented, changed even throughout the duration of this project, and it was not clear which APIs are compatible with each other. Ultimately the only reliable way to understand a class or function was to read its source code. In many cases these APIs provided almost what was needed, but to support our use case the code was copied from the TensorFlow GitHub repository and adapted to our purpose⁶.

The point of entry to the trainer program is the `controller.py`, which decides based on command line arguments which task to run: train, predict, train and predict, evaluate, export model, etc. There is a large number of hyperparameters that can be supplied via command-line arguments, with reasonable defaults.

The data was loaded using the `tf.data.Dataset` API, which provide convenient functions for reading large numbers of files, prefetching, shuffling, batching, and performing arbitrary transformations on the data, such as transforming a set of bytes representing a JPEG image into a 3D tensor of integers. This works well for simple use cases but is less flexible when for example doing multi-objective learning. We would like to control roughly how many data points from each objective end up in a batch, or for how long a model is pre-trained on one objective before the second objective is introduced. Refer to section 4.3 for a description of two approaches that were tried. In general, the trainer program would use `tf.data.Dataset` class to build an *input function* for either training or evaluating, and the input function would be supply data points for the train loop.

TensorFlow models were built using the `tf.estimator.Estimator` class by providing a custom *model function*. The model function would dynamically build the model based on the *model type* (deep / linear), *input type* (these are listed in section 4.2), *training objectives* and *hyperparameters*. An input type determines which input features are used and how they are represented, while a training objective determined which training dataset, loss function and evaluation metrics were used. Both input types and training objectives had simple configuration files dictating

⁵The author has experimented with a very similar set-up; latency was low enough..

⁶For example, multi-objective learning where each objective potentially changes a subset of the model's variables was not possible due to the way loss from multiple "heads" was merged in the current TF APIs. Also some parts had to be rewritten to give us per-class evaluation metrics.

3.2. SYSTEM ARCHITECTURE

their behaviour, which made adding new training objectives and experimenting with different model architectures easy.

Training was handled by the `tf.estimator.train_and_evaluate` function. It loads a model from the checkpoint directory if present, and trains the model for a specified number of steps, or until the input function terminates. It also periodically saves model checkpoints to GCS, and handles writing TF summary operations that are needed for data visualisations using TensorBoard

4. Method

Many kinds of experiments were conducted; the experiments, their evaluation and results are explained under the same subsection. In section 4.1 three visual similarity approaches are described and subjectively evaluated; in section 4.2 we describe all the models that were trained to reproduce the behaviour of the rule-based system; how multi-objective learning was used to improve the accuracy of individual training objectives is described in section 4.3; finally, our efforts to reduce label complexity using active learning is described in section 4.5.

4.1 Visual Similarity

Three approaches were tried for computing visual similarity. In the first case, an approximate nearest neighbour method [4] was used on the image embeddings (as described in section 3.2.2) and the top 100 predictions were saved to ElasticSearch. In the second approach, the embedding vectors were discretised according to [55] (as described in section 2.2.4) and inserted to ElasticSearch as strings of space-delimited tokens; standard ElasticSearch fulltext search was used on these token strings to get the nearest neighbour for a given product. The third approach combined the similarity scores of the second approach and the original ElasticSearch title matching, as a single ElasticSearch query.

4.1.1 Evaluation & Results

The evaluation of similarity scores is difficult, since it is inherently somewhat subjective. A common approach is to hand-label triplets of images Q, A, B with four choices of how similar A and B are to Q. These labeled triplets can be ranked using with similarity precision (percentage of triplets correctly ranked) and score at top K (see [68]). Creating this dataset of triplets would be a very time-consuming process. Therefore evaluation was done entirely subjectively.

Comparing the similarity between the former ElasticSearch title (ES title) matching and the new approximate nearest neighbour (ANN) search was straightforward - one look at the results was enough to confirm that the embedding-based approach outperforms title matching. Figures 4.1 and 4.2 show the nearest neighbours ac-

cording to ES title and ANN, respectively; the image embeddings nicely pick up the pattern on the shirt, as well as the general style. For some clothing categories this worked remarkably well, but with categories that had more varied products within it, there were occasional odd results. For example, in the “Hiking” category, the nearest neighbours of sleeping bags could be items with seemingly similar shapes, such as a flashlight with a very zoomed-in product photo that had similar contours as an unrolled sleeping bag. Another example of this is given in figure 4.4, which mostly returns chairs that are indeed visually and even functionally very similar to the chair in question - more so than the ElasticSearch title match seen in 4.3 - but also returns tables with similar legs; this is understandable, given that nearly half of the surface area of these images is covered by these foldable legs. This problem can mostly be mitigated by having more fine-grained product categories and restricting the nearest neighbour search to those more specific categories. The nearest neighbour similarity search may pick up qualities of the image that are not related to the product in question, such as background or fashion model; this is un-intended but is acceptable, and arguably creates a product listing which is more homogenous and therefore more aesthetically pleasing.

The second approach of discretised and tokenised image embeddings had mixed results. The returned nearest neighbours were reasonable, but subjectively felt inferior to the ones computed from the raw image embeddings using ANN. At times it felt like some aspect of the appearance dominated, which is understandable considering a lot of precision is lost when discretising the embeddings, and there is no theoretical justification why a TF-IDF score of these arbitrarily discretised tokens should behave similarly as a dot product between the original vectors. With enough precision (using several discrete intervals per embedding dimension) this may well work, but this would incur a considerable performance overhead when indexing and querying.

The third approach of combining the existing title matching and discretised and tokenised vectors was somewhat successful. As seen in figure 4.5, the title matching ensures that top results are at least about the same kind of items (chairs) while the tokenised embedding query returns products that share some visual similarity (mostly the legs of the chair). This approach still suffered from similar problems as the second approach - poor performance and somewhat inconsistent visual similarity - while introducing an additional complication of finding a good weighting between the title and visual scores. Given that there was no efficient way to evaluate the different weightings, it was decided to use raw image embeddings for visual similarity, either as an ANN search or a precise realtime version.

4.2 Independent Models

Several models were trained on the rule-based labels to find out which models can best replicate its behaviour. We first describe the Wide & Deep model that was trained as a baseline. Then the different input features and how these were

4.2. INDEPENDENT MODELS

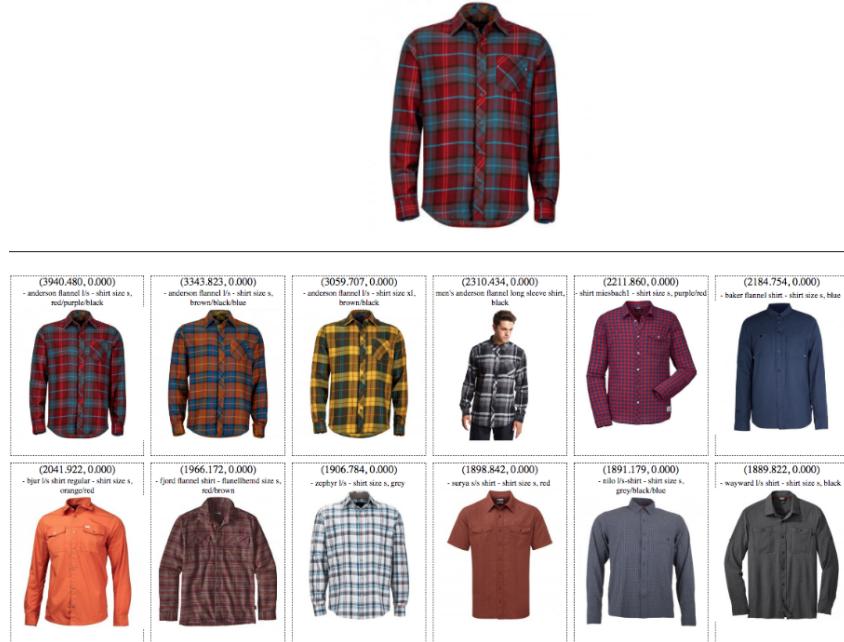


Figure 4.1: Nearest neighbours based on ElasticSearch title matching (TF-IDF-like)

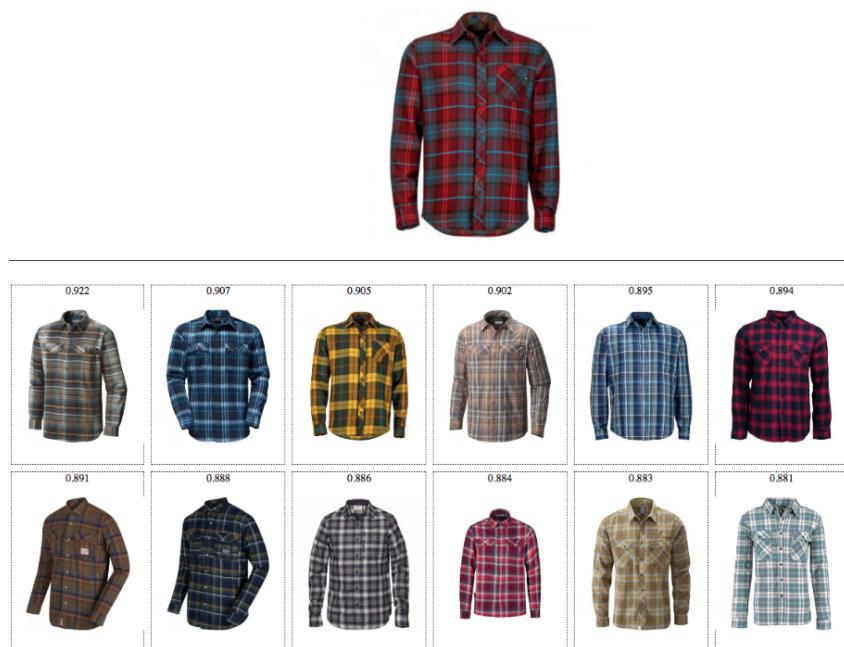


Figure 4.2: Nearest neighbours based on ANN search of image embeddings

CHAPTER 4. METHOD

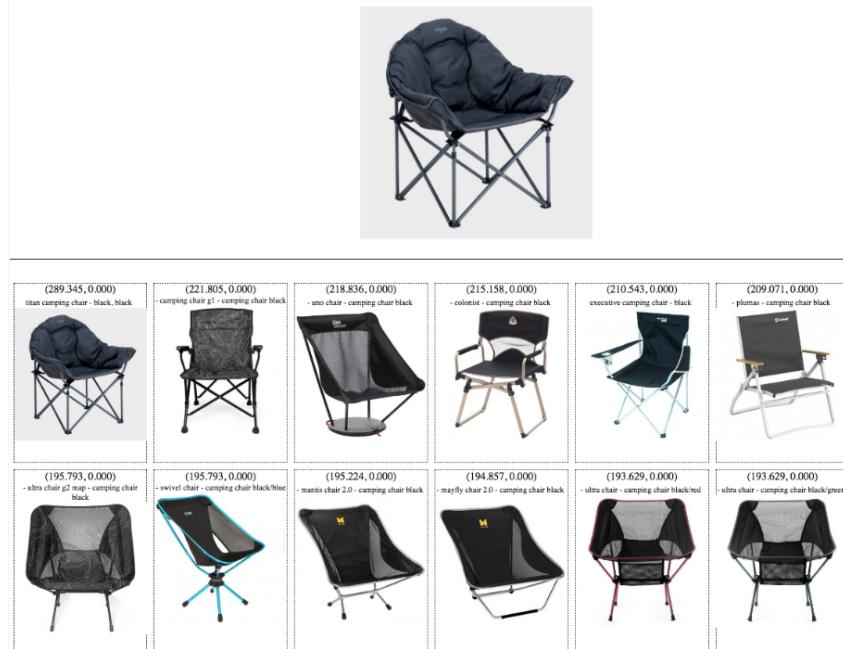


Figure 4.3: Nearest neighbours based on ElasticSearch title matching (TF-IDF-like)

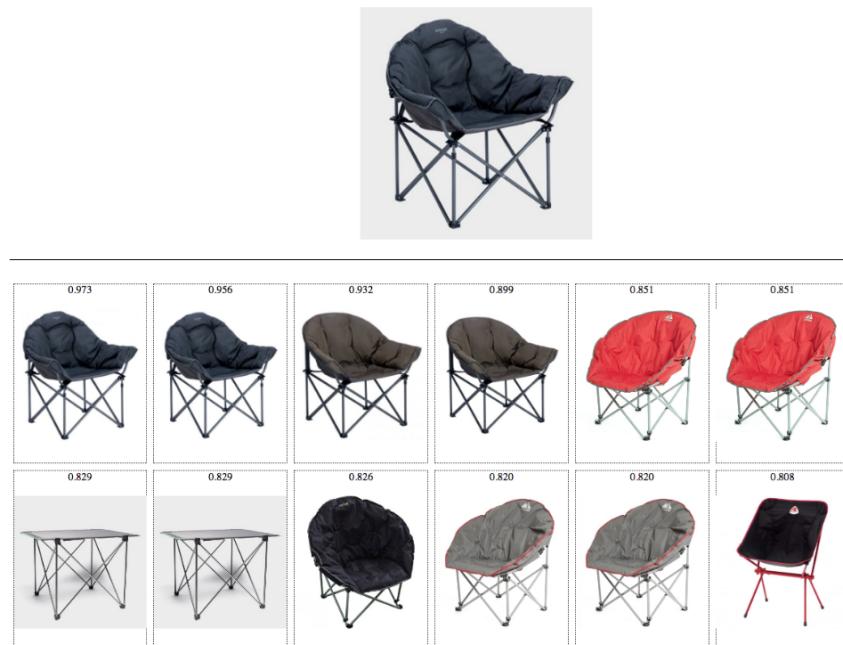


Figure 4.4: Nearest neighbours based on ANN search of image embeddings

4.2. INDEPENDENT MODELS

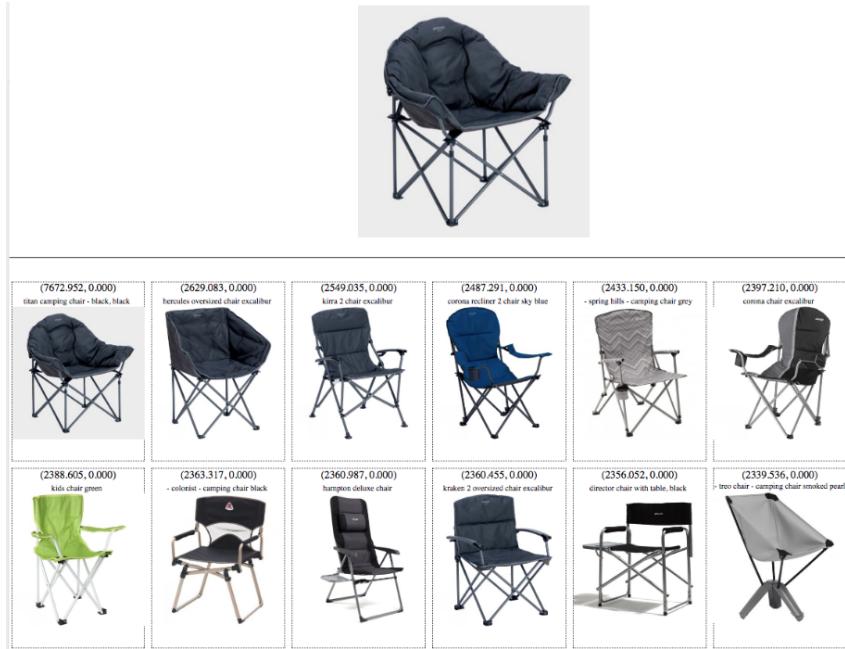


Figure 4.5: Nearest neighbours based on fulltext search on discretised/tokenised image embeddings as in [55]. Here scores from the title and image embedding fields was combined.

represented are described, along with how linear and deep models were trained from these input representations.

4.2.1 Baseline Model: Wide & Deep

As described in section 2.2.1, the “Wide & Deep” model consists of two models that are trained jointly using stochastic gradient descent: a deep and a shallow neural network which both predict the same set of binary outputs. The original idea of Wide & Deep was to use the wide component for feature crosses (e.g. of two categorical values), but in our experiments the wide component received mostly just the 1-hot encoded categorical inputs while the wide component received the same inputs as random embeddings (or as pre-computed image embeddings extracted with Inception V3 that is pre-trained on ImageNet).

Initial experiments with this model were done on a dataset of roughly 800 000 products which were labeled by the rule-based system¹. The data was partitioned into a train (~80%), validation (~10%) and test (~10%) sets in the Dataflow preprocessing pipeline based on a pseudorandom number generator - assign the product in question to the dataset if the random number is in the range 0 ... 0.8, 0.8 ... 0.9, 0.9

¹The database in the client company was in constant flux, so the training set sizes depended on how many labels the rule based-system had assigned to the current database. Still, when comparing different model architectures, we used a “frozen” version of the data to avoid skewing the results.

... 1. The train set was used to train the model, and the validation set was used to evaluate the model periodically as it was trained - during individual training runs and during hyperparameter tuning. The test set was held out for evaluating the model after hyperparameter tuning, but was not used due to the low variance in validation set error during the different tuning rounds.

4.2.2 Input Representation & Model Type Combinations

To assess which features are most useful for learning the rule-based objective, different features were grouped into *input types*, shown in table 4.6. The columns represent the input feature; refer to section 3.1 for the full feature names and their description. Entries in the table show how each input feature was represented: *emb* for categorical features that are represented as random embedding vectors (which are updated during SGD), *emb avg* for multi-token fields where tokens are represented as random embeddings and the input feature is a weighted sum of these embeddings². Feature columns marked with *u.s.enc* have been converted from raw text into dense embedding vectors of size 512 using a pre-trained model Universal Sentence Encoder [8]; cells marked with *mobilenet* and *inceptionv3* have respectively extracted image embeddings of size 1280 and 2048 using the Mobilenet [29] and InceptionV3 [63] models that were pre-trained on the ImageNet challenge; these embedding models were used via TensorFlow Hub³. Pre-trained models were used only as feature extractors, i.e. were not fine-tuned on our data.

Two model types were used to train with each of these 16 input types: linear and deep. Like in the case of the baseline model, the ~1.2 million data points were split into train/validation/test sets (80%/10%/10%); train and validation sets were used during individual training runs and hyperparameter tuning. All 32 combinations of input and model type were trained using hyperparameters listed in table 4.20 column *Deep (rb all)*.

4.2.3 Evaluation

Many metrics can be used to evaluate **multi-label classification**:

- TP, TN, FP, FN: true positives, true negatives, false positives, false negatives,
- accuracy: $\frac{TP+TN}{FP+FN}$, fraction of correct predictions,
- precision: $\frac{TP}{TP+FP}$, fraction of positive predictions that were correct,
- recall: $\frac{TP}{TP+FN}$, fraction of positive labels that were correctly “recalled”,

²Each embedding vector is weighted by the number of times it appears in the input feature, and divided by the L2 norm of the token counts; this gives good results for bag-of-words inputs according to TensorFlow documentation:

https://www.tensorflow.org/api_docs/python/tf/feature_column/embedding_column

³TensorFlow Hub is a collection of TensorFlow models that are pre-trained on some other task and can be used as a feature extractor or as an initialisation for a model that is fine-tuned: <https://www.tensorflow.org/hub/>

4.2. INDEPENDENT MODELS

Input Type	cat	raw cat	title	desc	brand	size	gender	img
full_emb	emb	emb avg	emb avg	emb avg	emb	emb	1-hot	mobilenet
full_one_hot	1-hot	k-hot	k-hot	k-hot	1-hot	k-hot	1-hot	mobilenet
noimg_emb	emb	emb avg	emb avg		emb	emb	1-hot	
noimg_one_hot	1-hot	k-hot	k-hot		1-hot	k-hot	1-hot	
categorical_emb	emb				emb	emb	1-hot	
categorical_one_hot	1-hot				1-hot	k-hot	1-hot	
raw_cat_emb_avg		emb avg						
raw_cat_k_hot			k-hot					
title_emb_avg				emb avg				
title_k_hot					k-hot			
desc_emb_avg					emb avg			
desc_k_hot						k-hot		
title_uni_sent_enc_emb						u.s.enc		
desc_uni_sent_enc_emb						u.s.enc		
mobilenet_emb								mobilenet
inceptionv3_emb								inceptionv3

Figure 4.6: Input types. Orange: embedding representation, blue: 1-hot or k-hot, red: embedding extracted with a pretrained deep network.

- F1: $2 \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}$, harmonic average of precision and recall,
- TPR = recall: $\frac{TP}{TP+FN}$, true positive rate, probability of detection,
- FPR: $\frac{FP}{FP+TN}$, false positive rate, probability of false alarm,
- ROC AUC: the area under the curve of TPR plotted against FPR.
- PR AUC: the area under the curve of precision plotted against recall.

In our case, an overwhelming number of labels are negative, therefore a very simple predictor (that predict “negative” for every product / category) would have an accuracy near 99%; ROC AUC is not suitable for the same reason, as a big change in FPs would result in a small change in FPR due to large TNs squashing FPR to a small number (see [13] for details). Precision and recall are useful measures to understand the model: it is expected to see a sharp increase of precision and decrease of recall as the model learns to predict “negative” for most classes, followed by a gradual increase of recall as the model learns to predict “positive” for the correct classes. Precision and recall are competing metrics: increased precision tends to lead to decreased recall and vice versa, therefore models which combine the two are suitable for evaluating multi-label prediction problems with unbalanced classes. While F1 score gives the harmonic mean of precision and recall at a given threshold, PR AUC paints a more complete picture by showing what the ratio would be at all the threshold values⁴. Therefore, all evaluation of multi-label classification is based

⁴There are infinitely many threshold values in the range 0 ... 1, so naturally we consider a small number of threshold values when approximating AUC.

on PR AUC, i.e. when choosing between model architectures or hyperparameters, a model with a higher PR AUC is preferred. Refer to [13] for an explanation why PR AUC should be a preferred metric for highly imbalanced binary classification problems.

4.2.4 Results

Wide & Deep The PR AUC on the 800,000-item dataset for this model after hyperparameter tuning was 0.9981 (see section 4.4 for a full description of the set-up). A manual examination of the predictions revealed that the top predictions for most categories was high-quality, but there was a large number of products below the decision boundary that actually belonged above it. A later run of the model on 1.2M products gave an PR AUC of 0.8647; this was using the hyperparameters that the tuning step produced, so it is likely that the decreased AUC PR score is caused by a bug⁵. There were occasional odd results from the model once it was trained on the 1.2M dataset; it was hard to tell which part of the model was causing it, or even whether all the input features are actually improving predictive power; therefore, a more systematic approach was tried, explained in the next section.

Input Representation & Model Type Combinations Table 4.7 shows the precise PR AUC and recall scores of the final time steps. The training loss function of linear and deep versions of noimg_emb is shown in figure 4.8. The PR AUC on the validation set is shown in figure 4.9, where each model is trained for 300 000 batches, which is roughly 6 epochs; similarly, figure 4.10 shows how recall changes during the course of training. The PR AUC on the final step (300k, or 220k for inceptionv3) is shown in figure 4.11; recall plot of the final time step was omitted, as it was very similar to the final PR AUC scores.

4.3 Multi-Objective Training

4.3.1 Category Structure

At the time of writing, there were roughly 1300 categories defined in the client database. Categories were structured in a way that is typical of e-commerce: categories can have child categories, which in turn can have child categories, etc.

Categories can be considered as independent (multi-label, sigmoid activation) or mutually exclusive (multi-class, softmax output layer); the difference is shown on a Venn diagram in figure 4.12. The common way to handle this is to assume categories are mutually exclusive. With exclusive categories, assigning a label to a

⁵Two explanations seem most likely: (1) the initial PR AUC was artificially high as some of the test set products ended up in the training set; (2) when migrating from the 800k dataset to the 1.2M dataset, the image embeddings were transferred over as an attempted optimisation approach, which may have mixed up image embeddings between products.

4.3. MULTI-OBJECTIVE TRAINING

	PR AUC (deep)	PR AUC (linear)	Recall (deep)	Recall (linear)
full_emb	0.9848	0.6974	0.9351	0.4613
full_one_hot	0.975	-	0.9037	-
noimg_emb	0.9805	0.5354	0.9243	0.3018
noimg_one_hot	0.9627	0.7169	0.8765	0.4492
categorical_emb	0.799	0.4936	0.6155	0.2638
categorical_one_hot	0.7748	0.5445	0.5919	0.3069
raw_cat_emb_avg	0.8808	0.4514	0.7231	0.2637
raw_cat_k_hot	0.8706	0.6211	0.705	0.358
title_emb_avg	0.9298	0.4195	0.8242	0.2388
title_k_hot	0.9215	0.5363	0.7904	0.291
desc_emb_avg	0.8388	0.417	0.6714	0.2286
desc_k_hot	0.8975	0.6688	0.7562	0.3932
title_uni_sent_enc_emb	0.749	0.4885	0.555	0.2808
desc_uni_sent_enc_emb	0.7337	-	0.5304	0.4305
mobilenet	0.7233	0.6479	0.5065	0.4152
inceptionv3 (220k steps)	0.6697	-	0.473	-

Figure 4.7: Evaluation metrics of the rule-based training objective on the final time step. Minus denotes a train run that was not scheduled (inadvertently).

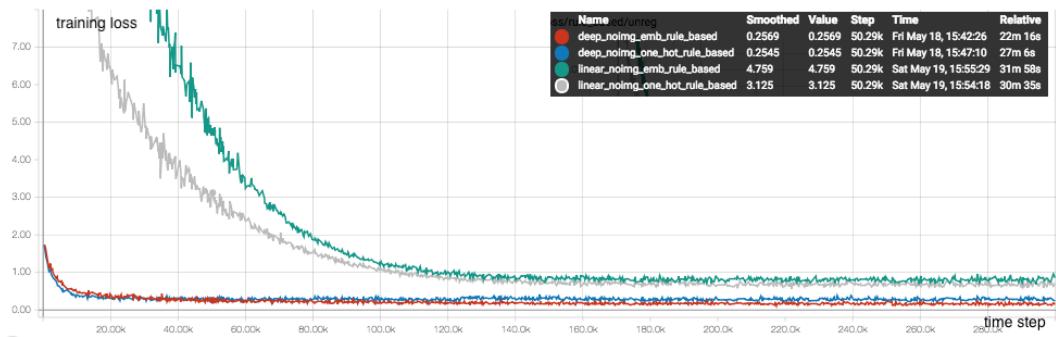


Figure 4.8: x axis = train step (batch number), y axis = training loss of rule-based objective.

CHAPTER 4. METHOD

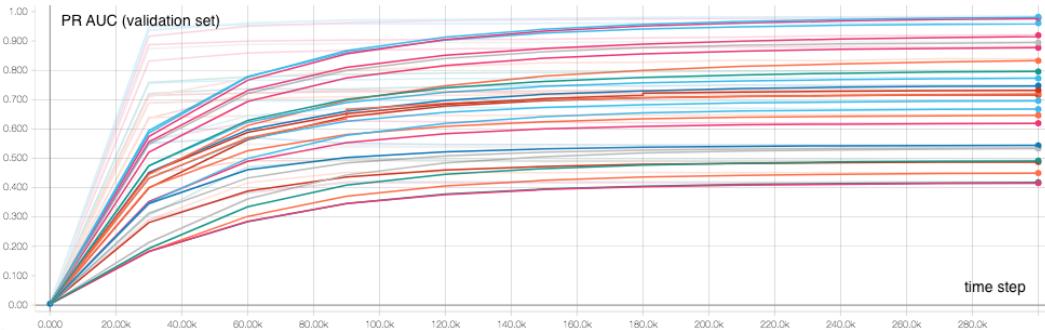


Figure 4.9: PR AUC of all combinations of input and model types; x axis = train step (batch number), y axis = PR AUC. Refer to figure ... for exact scores.

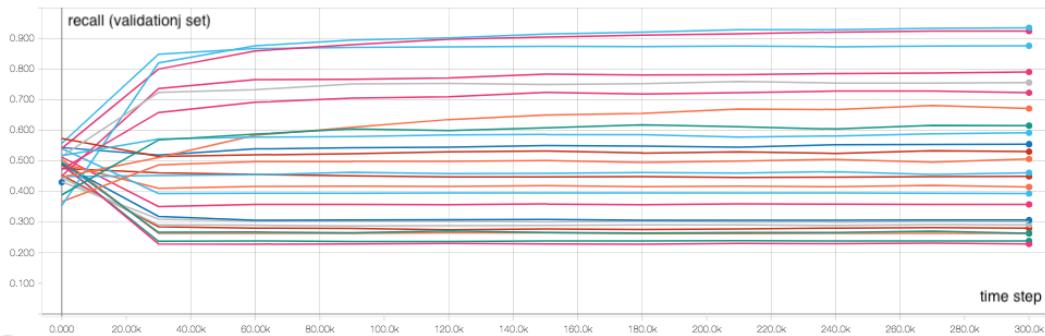


Figure 4.10: Recall of all combinations of input and model types; x axis = train step (batch number), y axis = recall..

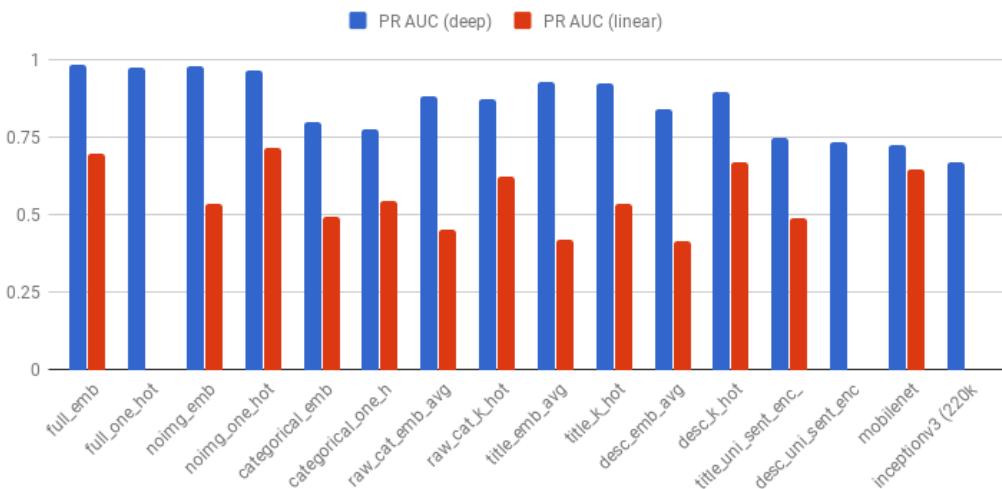


Figure 4.11: PR AUC on final timestep (300k, or 220k for inceptionv3)

4.3. MULTI-OBJECTIVE TRAINING

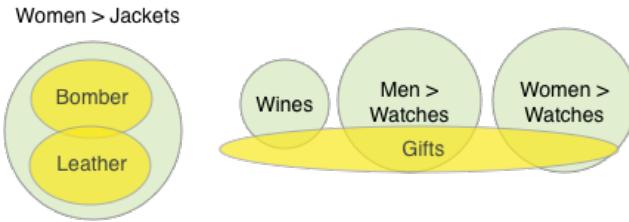


Figure 4.12: Exclusive (green) vs independent (yellow) categories. A product belonging to a mutually exclusive category would not belong to any other category (e.g. a jacket is never a wine), while a product can belong to many independent categories (e.g. the same jacket can be a bomber jacket as well as a leather jacket).

product determines its label across all categories, whereas with independent categories a positive label only determines the label regarding the category in question (and its ancestors in the category tree) - that is roughly a thousand-fold difference in labelling efficiency. The prediction task is also easier for exclusive classes - rather than needing to predict a probability score above a threshold separately for each class, the model would need to just assign the highest probability to the correct class compared to other classes. On the other hand, some categories at the client company are inherently ambiguous or even overlapping. The rule-based labels are also independent, since a product may be labelled to belong to zero or many categories.

Bearing in mind the considerable efficiency gain in labelling, we mark a subset of the categories (1063) as “exclusive”. Exclusive categories are usually leaves, but at times a higher-level category was marked as exclusive, denoting the more general case⁶. Therefore we have **two training objectives**: a multi-label prediction of independent categories labelled by the rule-based system, and a multi-class prediction of exclusive categories labelled by the employees of the client company - respectively called the **rule-based** and **exclusive** training objectives.

We can not use the rule-based labels directly for training the exclusive objective; even if they share the same set of categories, the logits (the value before the softmax/sigmoid activation) of multi-class and multi-label classification behave differently, so we can not just copy the parameters learned from rule-based objective to the exclusive objective. We can either do multi-objective training where the neural network has two output layers: one for the exclusive and one for rule-based objective (see figure 4.13 for a depiction of this multi-objective model). Alternatively, we could convert the labels assigned by the rule-based system (which are inherently independent) into exclusive labels; we call this training objective “**exclusivised**”⁷. This would allow us to train with only a single training objective, where most labels are assigned using the rule-based system and “exclusivised”, and a small number of

⁶e.g. if “Books”, “Books > Travel Books” and “Books > Cookbooks” are all marked as exclusive, then the semantics of “Books” is actually “Books that are not about travel or cooking”. This means some categories are not strictly exclusive, but we had to work around the existing category structure.

⁷This word does not exist in dictionaries, but its meaning should be clear.

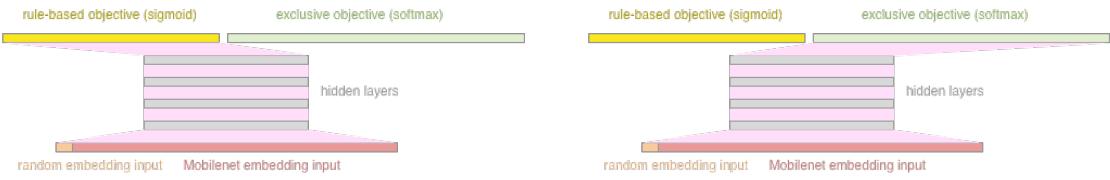


Figure 4.13: A multi-objective model. Left: gradients (pink) when a rule-based label is present. Right: gradients (pink) when an exclusive label is present.

labels comes from human labellers; this approach might have benefits as the output layer would receive far more examples compared to the exclusive objective in the multi-objective training scenario. To do that, we look at all the rule-based labels of a product that are applied to categories marked as exclusive. If there is only one such label assigned, as would be for most products, the label is already exclusive; if there are more than one such label, we can either insert the product to the training set twice (once for each label), or select a random label. We chose the latter approach.

4.3.2 Labelling of Exclusive Products

There were three purposes for labelling products: (1) creating a validation set, (2) creating an initial dataset for training the exclusive outputs so that reasonable uncertainty scores could be computed for active learning, and (3) labelling during active labelling rounds. There are three ways of picking products to be labelled: random sampling (a product is sampled randomly from all products that do not have an exclusive label), uncertainty sampling (the unlabelled product with the highest prediction entropy is chosen), and searching for a product manually using our web UI. When saving a label, the relevant metadata about how the product was sampled is also persisted to enable filtering out certain types of labels when running experiments.

In the experiments of the following section, the validation set contained 509 randomly sampled products and 5439 products that were manually searched for; for training set the equivalent numbers were 509 and 5646. Initially random sampling was used to obtain labels for products, to avoid bias in gathering data only through title search. Yet this approach would have taken a long time to populate all categories in the validation set with products, therefore the labelling team went through each existing category one by one and added 5 products to the test and validation set using title matching.

4.3.3 Evaluation

As described in section 4.3.1 we have three distinct training objectives: *rule-based*, *exclusive*, and a hybrid *exclusivised*. The models described in the previous section were trained only on the rule-based objective. In this section we see how a per-

4.3. MULTI-OBJECTIVE TRAINING

formant models from the previous section performs when trained on the exclusive objective (from a small number of hand-assigned labels), how jointly training using both the rule-based and exclusive objective improves the accuracy of the exclusive objective, and how it compares to the “exclusivised” training objective.

For the exclusive objective, which is an instance of **multi-class classification**, accuracy is a reasonable measure. Even though classes are imbalanced (“Dresses” has more products than “Travel Books”), there is no single dominant class that could be used as a shortcut by the model. On the other hand, similar problems arise with evaluating multi-class problems where there is potentially overlap between some classes: the metric would give a low score even if the model predicted a semantically very similar class, or a more general class instead of the more specific one it was labelled as. Therefore a manual evaluation of the misclassified products is needed to determine whether the error rate is caused by the ambiguity of the classes or by problems with learning.

4.3.4 Experiments & Results

Combining Data From Objectives

To train jointly for the rule-based and exclusive objectives, the data from the different datasets had to be combined. In the naive version, data from both datasets is loaded into memory and shuffled, but this would not work when we are also loading images or the dataset size grows. As a scaleable workaround, we read the file names of each objective into memory (a few hundred to a few thousand in our case), shuffle the file names, and alternate between reading m files from the first objective and n files from the second objective (we used $m = n = 1$); when one dataset runs out of files, wrap around, but if the other objective also runs out of files, terminate producing filenames. The resulting stream of filenames is read by a `tf.data.Dataset`, which loops through the produced filenames indefinitely and keeps a shuffle buffer of b elements; within this buffer the data points from the different objectives get mixed up, so each training batch usually gets products from both objectives. Note that the file shards are of different sizes: how many products end up in each file depends on how many Dataflow workers were used and how load was divided among those workers. Therefore, to avoid long periods where batches are filled with data points from a single objective, larger shuffle buffer sizes should be used.

Another simple option for combining data from different objectives is to dump all data in the same file, and optimise the model’s parameters with respect to all the objectives for which your items in the given batch have labels for. This ensures that data points are seen an equal number of times, but is problematic when the different objectives have vastly different number of labels, as there is no good way to force training on the under-represented objective more heavily. An additional unforeseen problem occurred with this set-up: due to the way data was pre-processed, the data points from the exclusive objective often ended up in a single file, which resulted in terribly unstable training unless the shuffle buffer was the size of the whole dataset.

This can be seen in figure 4.14, where the training loss of the rule-based objective (magenta) shoots through the roof once it encounters its first exclusive data point. In this case it happened a late stage, when the model had trained for nearly 4 epochs' worth of data⁸, and by that time the adaptive learning rate of Adam has already learned that the weights of the exclusive output layer are never updated, i.e. the next update should have a very strong learning rate. Now when the model sees all the exclusive labels at once, the gradients are big enough to affect the whole network, which is why the changes are visible even in the loss function of the rule-based objective; many other metrics (such as L2 regularisation loss and gradient norm) also grew sharply at the time step when exclusive and rule-based losses blew up.

Pretraining on the Rule-Based Objective

The initial plan was to pre-train the model on the rule-based objective, and then train jointly on both objectives. As seen on figure 4.14, the training loss of the rule-based objective (blue) makes large periodical jumps, followed by periods of high-variance downwards fluctuation. The first such jump starts after the first epoch of pre-training, and following jumps can leave the loss function plateaued on the level it jumped to. The sharp jumps can be explained by the high learning rate set for the parameters of the exclusive output layer that have received few or no updates, and the subsequent plateau of the loss function could be explained by the fact that the learning rate for the rest of the model has been reduced substantially, and the model (which is shaken up from its position near an optimum) is unable to escape its new sub-optimal “valley” of the loss surface. Contrast this with figure 4.15, where the jointly trained model's loss (red) is gradually converging towards the loss of the single-objective loss (blue). In figure 4.14 all the other parameters are the same as in 4.15; the former uses pre-training and the latter does not, and obviously the runs had different outcomes when shuffling data.

Exclusive, Exclusive & Rule-Based, and Exclusivised Objectives

Our real goal is to properly predict the exclusive objective. As is evident from figure 4.16, the model trained with just the exclusive objective converges fast and improves little after the first epoch, while the stable multi-objective model continues to improve in accuracy after several epochs, surpassing the single-objective accuracy after about three epochs. The PR AUC of a model trained on just the rule-based objective exceeds that of the multi-objective training, which can be explained by the additional variance the exclusive objective introduces. Both results are as expected, though the overall accuracy of 60% is not particularly strong. A reasonable guess

⁸Due to the way the input function that feeds data to the model is restarted every time the model is evaluated, there is no guarantee that all files/data points are accessed the same number of times. So it is possible the chunk with all the exclusive labels comes at a late stage or several times.

4.3. MULTI-OBJECTIVE TRAINING



Figure 4.14: Unregularised training loss of rule-based objective. Red: only rule-based objective. Blue: until step 35K, only rule-based objective, then a combination of rule-based and exclusive. Magenta: rule-based and exclusive objectives loaded from the same files. Shuffle buffer size: 15K.

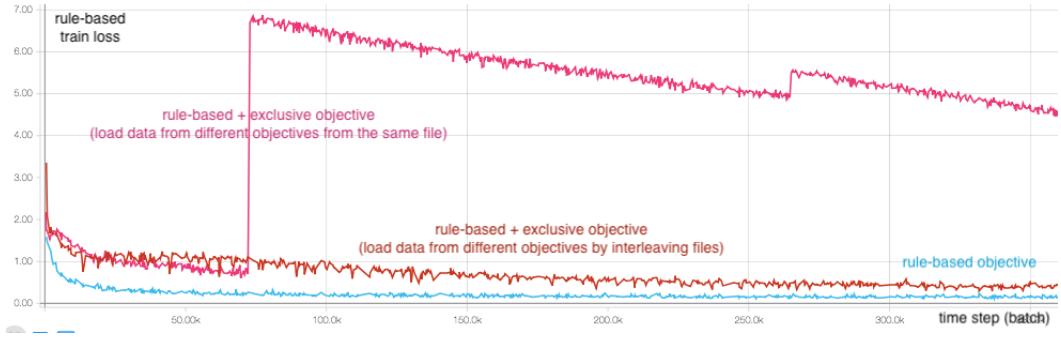


Figure 4.15: Unregularised training loss of rule-based objective. Blue: only rule-based objective. Red: combination of rule-based and exclusive objective. Magenta: rule-based and exclusive objectives loaded from the same files. Shuffle buffer size: 15K.

is that the parameters of the exclusive output layer did not receive enough training examples to learn from, so a model was trained on the exclusivised data set as well. Alas, as seen in figure 4.17, the exclusivised training objective does not outperform the former multi-objective learning and has a very high variance in the training error; a larger shuffle queue seems to help both settings, however. As training loss that has a high learning rate would fluctuate similarly when the learning rate is too high, a hyperparameter tuning job was scheduled to find out whether a lower learning rate would provide better results. Alas, this was not the case, and the poor performance of the exclusivised objective is most likely caused by the noise in the labels.

To assess whether the low accuracy of 60% is caused by a genuine problem with learning or an ambiguous validation set, a confusion matrix was computed on the validation set. Considering there were 1000 categories, the heat map of

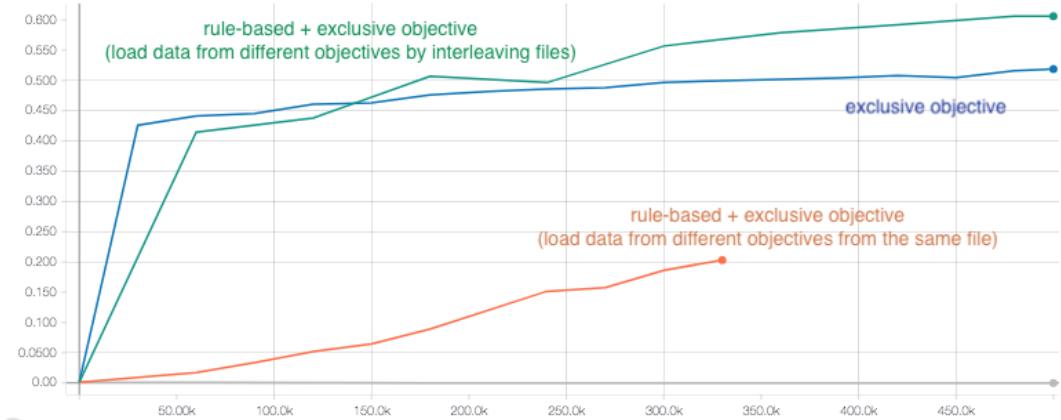


Figure 4.16: Validation accuracy per time step. Green: multi-objective model (exclusive, rule-based) with interleaved input files. Blue: only exclusive objective. Orange: multi-objective model (exclusive, rule-based) loaded from the same training file (unstable). Shuffle queue: 15K.

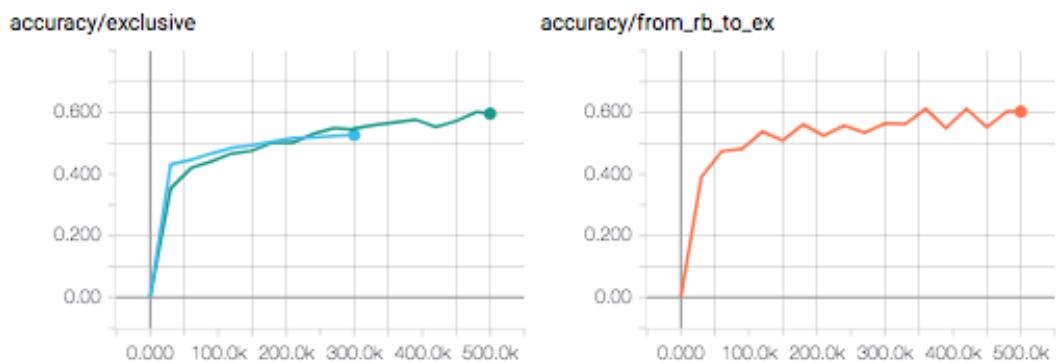


Figure 4.17: Validation accuracy per time step. Left green: exclusive objective in multi-objective training. Left blue: exclusive objective. Right orange: exclusivised objective. Shuffle queue: 100K.

4.4. HYPERPARAMETER SELECTION & TUNING

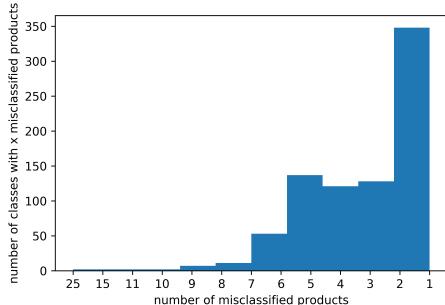


Figure 4.18: Number of misclassified products in validation set.

class	misclassified	class	misclassified
Knee Length	25	Casual Shoes	9
Games	15	Clothing	9
Hockey	11	Heating & Coolin	9
Kids Luggage	11	Perfumes	9
Telescopes	10	Long-sleeve Tee	9
Dental	10	Trousers	8
Mini	9	Shoulder bags	8
Casual Shoes	9	Parts & Tools	8
Mens Clothing	9	Tall	8
Accessories	9	Travel Tech	8

Figure 4.19: Worst performers in validation set.

the confusion matrix was too dense to reveal useful patterns. Instead, a histogram of the number of categories that have a given misclassification rate was computed (figure 4.18); note that this was computed from both sides of the diagonal of the confusion matrix, therefore a misclassification is counted twice. Top 20 categories with the highest misclassification rates are shown in figure 4.19. The categories with many misclassifications tend to be very general ones; for example, Clothing, Men’s Clothing and Casual Shoes are certainly not exclusive, and the client company is advised to find ways of expressing these as exclusive categories. For some (e.g. “Knee Length”, “Mini”, “Tall”), visual information could be useful which in this round was not used. The large number of categories that have a single or a handful of misclassifications is most likely caused by a learning problem such as insufficient labels or suboptimal training.

4.4 Hyperparameter Selection & Tuning

Hyperparameters were chosen based on gut feeling and Bayesian optimisation. The former was used to come up with reasonable ranges for hyperparameters; automatic tuning was done on the GCP ML Engine, where one just needs to specify the ranges of hyperparameters, their type (categorical, integer, reverse log scale), and how many models are trained in these hyperparameter ranges. The ML Engine tuning mechanism uses Bayesian optimisation to come up with candidate hyperparameter values, which is described in section 2.5.1 and by GCP engineers⁹. The hyperparameter ranges used and the results that each of the three tuning rounds found is shown in table 4.20. In all cases 60 models were trained, with maximum 4 models trained in parallel, with early stopping.

The *L2 scale* parameter is the weighting factor on the L2 regularisation term that is applied across the weights of the linear and deep models. *Linear learning*

⁹<https://cloud.google.com/blog/big-data/2017/08/hyperparameter-tuning-in-cloud-machine-learning-engine-using-bayesian-optimization>

	Wide & Deep		Deep (rb title)		Deep (exclusivised)		Deep (rb all)
	full range	found	full range	found	full range	found	used
batch size	32, 64, 128, 256	64	64	-	64	-	64
L2 scale	-	-	$10^{-1} \dots 10^{-5}$	0.099379	0.0005	-	0.0005
linear learning rate	10^{-3}	-	-	-	-	-	-
deep learning rate	10^{-3}	-	$10^{-2} \dots 10^{-5}$	0.002657	$10^{-4} \dots 10^{-6}$	9.99×10^{-5}	10^{-4}
dropout	0, 0.1, 0.25, 0.5	0	0, 0.1, 0.25, 0.5	0.5	-	-	0
layers	1 ... 30	17	1 ... 30	28	4	-	4
hidden units	30 ... 1000	784	30 ... 1000	85	784	-	784
activation	relu, tanh	relu	relu	-	relu	-	relu
	PR AUC: 0.998		PR AUC: 0.4069		Accuracy: 0.468		

Figure 4.20: Hyperparameter ranges and the values found during automatic tuning. A single value in the “full range” column means the value was fixed.

rate is the learning rate of the Adam optimiser for the linear part of Wide & Deep, or the weights of the linear model; the other parameters of the optimiser were left as defaults (in TensorFlow and the original paper [40]); similarly for *deep learning rate*. Dropout corresponds to the likelihood of dropping out a hidden unit during training; *layers* corresponds to the number of hidden layers, and *hidden units* corresponds to the number of units in those hidden layers; *activation* corresponds to the activation function used in the hidden layers; the four hyperparameters are only applicable to the deep and Wide & Deep models.

Table 4.20 shows the the tuning rounds in different colours, as well as the PR AUC achieved by each. In the *Wide & Deep* case, tuning was performed over ~800k data points on the rule-based objective, with each model trained up to 6 epochs (e.g. 125k time steps with a batch size of 64). In the *Deep (rb title)* case, a random sample of 10% of the full 1.2M dataset was used to train on the rule-based objective for up to 50k time steps; the model used was *deep_title_emb_avg*, which showed good results compared to the small selection of models that was manually tried. In the *Deep (exclusivised)* case, the *deep_noimg_emb* model was trained on the full 1.2M dataset on the *exclusivised* objective (described in section 4.3) for up to 300k time steps. In the *Wide & Deep* and *Deep (rb title)* cases, models with a higher PR AUC score were preferred while in *Deep (exclusivised)* accuracy was used to choose between hyperparameter performance.

It is clear that the hyperparameters found by the *Deep (exclusivised)* round were not optimal, therefore the ones used in the *Deep (rb all)* was used in most experiments described in section 4.2.2.

4.5 Active Learning

A limited experiment with active labelling was conducted. After obtaining a small set of exclusive labels (as described in section 4.3.2), the model *noimg_emb* model was trained in a multi-objective setting, and the entropy of the exclusive objective was used (averaged across all exclusive classes) to sample products. Products were

4.5. ACTIVE LEARNING

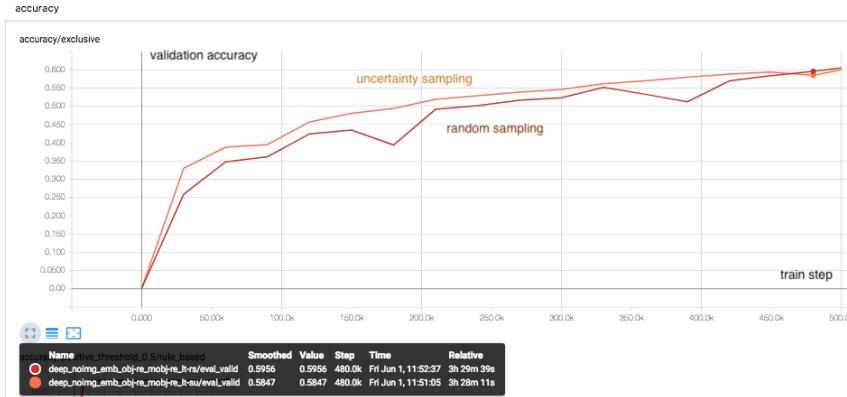


Figure 4.21: Active learning: validation accuracy of the exclusive objective when training on a small number of randomly sampled vs uncertainty-sampled products.

sampled 1-by-1 (deterministically, highest entropy first)¹⁰ and labelled by a human. The most confident prediction was pre-selected for labelling convenience.

The initial idea was to perform active labelling in rounds, where a the same number (e.g. 1000 - 5000) of products would be labelled via random and uncertainty sampling. After each round, three models would be trained: one model that uses all the labels, one that uses only uncertainty-sampled labels, and one that uses randomly sampled products. Labelling products 1-by-1 turned out to be impractical, which is why only a single round of 200 products was labelled (100/100 random/uncertainty sampled), which were added to the 6155 products in the training set of the exclusive objective.

4.5.1 Results

The products returned by this method were in some sense what one might expect, at times returning products that were tricky to label. At times the product was hard to label because there actually did not exist an appropriate category for it. During this brief labelling round, 12 new categories were created - in all cases when trying to label a product sampled based on uncertainty. One unfortunate side effect of this sampling strategy is that it tends to return products that are often similar to one another. The final metric of interest (validation set accuracy of the exclusive objective) is not changed much by the sampling strategy, however as shown in figure 4.21, uncertainty sampling seems to speed up convergence somewhat - though sample size is too small to make any definitive claims.

¹⁰Ideally we would have also tried disagreement sampling, but this would have involved loading predictions from different models across a large number of large files, which is not only slow but can easily introduce bugs related to misaligned predictions (i.e. products are in a different order across the predictions of different models). To store probabilities for 1000 classes across 4M products as 32bit floats, one needs 16GB; multiply that with the number of models you need to compare uncertainty between.

5. Discussion

In this section we analyse the experiments that were described in the previous section.

5.1 Visual Similarity

Although this assessment is subjective, the results from the approximate nearest neighbour were strikingly good; this is the opinion of many people at the client company that saw a demo of it. The model did not require fine-tuning, contrary to the initial hypothesis. The biggest challenge here is to find a way to evaluate visual similarity performance; although there are various algorithms for tweaking similarity and evaluating it on a test set, creating such a dataset would have been prohibitively expensive. Given that vanilla pre-trained CNNs produce good embeddings for visual similarity, it is sensible to focus one's efforts on the scalability and productionalisation aspect of this problem. Still, two relatively easy directions could yet be explored: the model architecture and the layer from which embeddings are extracted. Our model of choice was Inception V3, which had a good accuracy to memory footprint ratio - with respect to the original ImageNet challenge. However it has been shown that models that are well tuned for a given classification task are not necessarily the best feature extractors; in fact, ResNets are currently the best at this [43]. The other decision of using the penultimate layer of the CNN as a feature extractor is also a common one. Often the transfer learning task is classification, for which the higher-level fully-connected layers would provide good high-level semantic features. In our case, extracting features from the first fully-connected layer or even from the raw convolutional feature maps at a lower layer might have given us similarity that is highly sensitive to certain visual patterns and shapes - which could be deployed to certain categories where such sensitivity is required.

As already discussed, the tokenised versions of feature vectors had less spectacular results. While appealing as a way to combine visual and textual similarity, there is not theoretical justification for this approach. Note that the original paper [55] extracted dense features from text, and only hypothetically proposed this approach to be applied on images. Perhaps the inherent discrete nature of language (sentences, words, topics as opposed to a continuous 2D space) makes it more amenable to such tokenisation.

5.2 Individual Models & Hyperparameters

Choosing a relatively complex model such as Wide & Deep as a baseline model can be criticised: a baseline should be simple and interpretable. One of the justifications was the ability to turn some part of the model off, as removing all deep input columns would have resulted in a linear model and vice versa¹. The other intuition was drawn from models with skip connections such as ResNets [24] and DenseNets [31]. The assumption was that linear models are relatively good at predicting the output, as the relevant word tokens would be present in the textual fields of many products; for a deep model to use this raw knowledge, it would need to learn an identity function from the input of the layer to the output, which can be a hard problem with some activation functions. Therefore, a linear layer with the same input as the deep layer could be justified to simplify learning for such simple cases, though not necessarily as the first “baseline”. Overall Wide & Deep performed very well, though it is hard to compare its PR AUC to the other models, which were trained and evaluated on twice the amount of data. When its performance degraded once it was trained on a new dataset (that most likely had mixed up image embeddings of some products), its drawback became apparent: we had no clue which of the input features were most responsible for good or bad performance.

Training a large selection of deep and shallow models was a good way to get a sense of the capabilities of these models. The performance of linear models was lower than expected, but this is probably not a fair comparison, as the models were trained on the same sets of hyperparameters that were at least in part chosen to be suitable for deep models. In general the tools used (TensorFlow, Adam optimiser) may not be optimal for learning linear models, but considering the large difference of PR AUC between the two, there seem to be benefits to using deeper models. Contrary to our hypothesis, deep models with 1-hot / k-hot inputs outperform shallow ones substantially, but the difference in performance decreased when the models received a bigger selection of input features. This implies that deep models are indeed generalising better from a small set of inputs; for example, the deep model that received averaged word embeddings of the title continued to improve even at epoch 6, and received a remarkably better PR AUC score (0.93) than its linear counterpart (0.42, or 0.53 for linear with k-hot inputs).

The biggest problem in interpreting the PR AUC scores of the rule-based objective is that its labels were generated using a relatively simple automatic procedure which only considers text as its input. Therefore a high PR AUC score shows us only how well the model was able to reproduce the behaviour of the rule-based system, but we are after good generalisation and consistent multi-objective performance. Therefore good results from title-only models should be taken with a grain of salt. We consider the classification problem to be a quite easy optimisation problem, however the power of more complex models might be better served when we are adding more training objectives (e.g. colour, style, conversion of products).

¹Implementing separate model types (deep, linear, wide & deep) turned out to be simpler.

5.3. MULTI-OBJECTIVE TRAINING

Some patterns emerge from training this selection of models. For linear models, 1-hot and k-hot encoding always outperforms embedding inputs, which was contrary to our hypothesis. Deep models pre-trained on other objectives are quite consistent in their generalisation, producing strikingly similar PR AUC and recall scores whether it takes the title, description or image as input. Examining the true/false positives/negatives of individual classes reveals that these models are relatively good at predicting certain classes, yet fail to differentiate between more specific ones.

It would be interesting to compare the performance of these models with random forests and 1D CNNs; fine-tuning of pre-trained models would also be an interesting comparison, as would models trained on character n-gram representations or TF-IDF-weighted embeddings. Our embedding models could further be improved by not limiting the number of unique tokens so aggressively; as embedding based models have shown to perform better than 1-hot encoding in deep models, we could expand the vocabulary sizes without increasing input dimensionality.

Hyperparameter tuning using Bayesian optimisation was not as successful as expected. In the first round where the Wide & Deep model was tuned, the final metric value was only marginally higher than PR AUC of manually chosen hyperparameters, yet the model complexity was far greater. A simpler model with comparable accuracy is more likely to generalise better, therefore most of the following deep models used 4 layers. In the second tuning round, the model failed to learn in any instance; this is probably because the large range supplied for the “L2 scale” parameter. An examination of the individual models that were trained during tuning showed that L2 decay overpowered the training loss and resulted in a very strict model that nearly always predicted the negative class. The rate weight decay scaling parameter was kept fixed at 0.0005 for future runs based on analogous values used in literature; there is also some evidence that one should specify a custom scheme of weighting the L2 loss with adaptive optimisers such as Adam [46].

Given that we were training models with high capacity, it would also be useful to know the variance of the outcome if the model is trained several times by doing k-fold cross-validation. Due to the large dataset size and time limitation this was not tried formally, however these models were trained across dozens of different data pre-processing runs during which the train/validation split was different. There was not much variability within even the high-capacity models, which often converged at a high AUC PR score.

5.3 Multi-Objective Training

Even though the rule-based objective gave good PR AUC scores, the model was not deployable as is, since it severely under-predicted all nontrivial categories. As one might expect, training for multiple objectives has a higher variance in the gradients. This was exacerbated by the limitations of the tf.data APIs, which did not enable to efficiently alternate between the training objectives precisely. Probably the variability could be reduced if each batch had a consistent number of products

from each objective.

Although Adam was a versatile optimiser across the different deep and shallow architectures for the rule-based objective, updates to the parameters of each top layer becomes less frequent in the multi-objective case. As Adam learns to update less frequently updated variables more heavily, this is probably the cause of a wide range of issues described in section 4.3. Experimenting with a simpler optimiser and outputting per-step statistics about how many data points from each objective ended up in each batch would help further understand the shape of the loss functions.

It is satisfying that a multi-objective training scheme outperforms a method trained on just exclusive labels. Though the exclusivised objective at times exceeds multi-objective training in terms of accuracy, the latter is preferred due to its stability and predictability; also, this scheme will be useful once we add additional training objectives. The 60% on the exclusive objective is not ideal. This is achieved with roughly 6000 manual labels, which is far from the millions of labels deep neural networks usually require. Probably the biggest gains will come from re-organising the category tree further so that there are no ambiguous classes, and that all the ambiguous cases (e.g. “Leather Jackets”) are handled as a combination of an exclusive class (“Jacket”) and independent feature detector (“Material: leather”). The remaining gains will come from active labelling, which should find the uncertain cases, or by manually adding labels to categories with a high error rate. An orthogonal approach is to train using the Wasserstein loss, which would be better at handling ambiguous and noisy labels; in this case, we might use the Wasserstein distance between the predicted and actual labels as the evaluation metric as well, as accuracy would still paint a pessimistic picture when in reality it is acceptable to classify into a semantically nearby class. Different training regimes could still be tried to stabilise multi-objective learning, such as gradient clipping, layer and batch normalisation.

5.4 Active Learning

Preliminary results indicate that uncertainty sampling is effective at identifying products that might lack a category or sufficient labels in the training data, but it tends to sample products that are very much alike. Therefore this strategy may not be very well suited for such a “batch-active” learning scenario that was presented. A workaround might be a series of smaller training rounds. Sampling products 1-by-1 is perhaps the biggest disadvantage of this approach, which results in a time-consuming process. An alternative would be a system which presents the user with a page of predictions to confirm/reject, which would apply a large number of labels.

6. Conclusion

In this work we covered a large breadth of topics. Neural networks and deep models have many useful qualities: their performance continues to improve with additional training data, they can be successfully used for transfer and multi-task learning, and they can synthesise realistic samples of their own. In general, deep models based on embedding inputs outperform ones with sparse inputs, while linear models do well with 1-hot encoded inputs. Transfer learning is remarkably useful, as our experiments with visual similarity show, though models trained from scratch on raw product data is still superior for fine-grained affiliate product classification. While this work did not experiment with ensembling of models, we did look at joint training of models that individually performed reasonably well, and as was expected found that a combination of all these inputs outperforms any subset.

A big part of the time spent on this project involved building the technical architecture. This was a non-trivial task given the large number of technologies, distributed components and dataset sizes involved. This could serve as a blueprint for many projects that are planning to deploy to the GCP cloud.

Our experiments with multi-objective training ultimately produced the expected results, and the performance increase could be hopefully improved with more analysis of the high gradient variance during training. Even though a good number of models were trained to optimise hyperparameters using Bayesian optimisation, the bulk of these were still picked based on intuition. At the very least, such tuning reveals that the models are not too sensitive to hyperparameters in this problem.

The results from active learning are not yet conclusive, though the tentative conclusion is that uncertainty sampling should not be used too early when the model is not yet close to an optimal solution, as it might sample very similar data points. Further work and experiments with active labelling should be conducted to find out whether it outperforms other sampling approaches.

Overall, the choice of tools and approaches was reasonable. There are clearly more straightforward ways to do product classification, yet the complications introduced by transfer, multi-objective and active learning will most likely pay off when predicting much more difficult aspects, such as whether a product would sell well for a given audience or not - which has immense potential to increase profitability.

6.1 Future Work

The following is a summary of some of the possible directions for future work and research:

- Use a training scheme where each batch contains a consistent number of data points from each training objective. This has been implemented after the defence of the thesis¹, which removes the majority of training instability.
- Use bulk labelling of products. Present the user with a page of products that the model predicts to belong to a given class; sort the products either by the “confidence” of the prediction (output from the softmax layer), or the entropy of the prediction. The former case helps correct predictions that are confidently wrong, the latter helps the model learn from more challenging cases. Preliminary results show great promise in terms of labelling efficiency and increase in model accuracy.
- Organise the category tree such that there are no overlapping exclusive categories. In one scheme, categories would be considered either exclusive or independent; exclusive categories can have independent children, but not vice versa. Exclusive categories are considered as a multi-class classification problem as before, but independent categories are considered as multi-label classification - with the caveat that independent outputs are only predicted/optimised when their parent exclusive category is predicted/labelled. Preliminary results show accuracy gains from 60% to 80% on the exclusive objective, and good label efficiency on the independent outputs.
- Examine whether multi-objective training continues to improve the accuracy of the objective with fewer labels when the objective receives an increasing amount of labels. Multiple objectives might compete with one another unless the objectives are very similar.
- Experiment with optimisers that do not use adaptive learning rates.
- Experiment with Wasserstein loss to penalise misclassification to semantically nearby classes less than misclassification to semantically dissimilar classes.
- Experiment with training techniques that may further improve training: layer/batch normalisation, gradient clipping, and separate weighting schemes of the loss of different objectives.

6.2 Practical Take-Aways

First and foremost, this work was a practical exercise in reproducing many results found in literature. As such, there are some non-academic points that the author

¹There is a solution to this even using tf.data.Dataset:
<https://stackoverflow.com/questions/49981542/tensorflow-concat-tf-data-dataset-batches/49990928#49990928>

6.2. PRACTICAL TAKE-AWAYS

learned, and could be useful for other ML practitioners.

- TensorFlow high-level APIs enable writing very succinct code, but unless the use case is very simple, it is not sufficient to use canned Estimators. When in doubt what a function does, read its source code and add `tf.Print` and `print` statements to examine what the tensors passed around actually look like. When working on non-standard models, create your own Estimator and Head by seeing how canned versions were implemented. Using these high-level APIs will eventually pay off in terms of deployment ease and standardisation.
- Apache Airflow is a very popular tool for building data pipelines, probably currently the best solution. It is good for building pipelines where the inputs are known up front, but less suitable for rapid experimentation where one might want to more conveniently provide dynamic arguments to sub-tasks (there is a solution for that in Airflow, but it is somewhat awkward). The plug-ins for GCP in Airflow were not as feature-rich as one might expect, therefore it was often easier to create custom tasks that simply executed the relevant command-line `gcloud` commands. Additionally, at times the scheduler randomly did not schedule some tasks, and if the reason for it was logged somewhere, it got lost in the rest of Airflow’s verbose logging system.
- On a big dataset, handling images as the input to the pipeline/model adds several orders of magnitude in extra computation. In our case, the benefits were marginal, and had the only task been product classification, then including images as model inputs would have been impractical.
- Latencies in scheduling a Dataflow or ML Engine job are on the order of 5 min. That is too slow for experimenting on such deployment platforms alone. It is very beneficial to have a small dataset for experimentation, which can run the full pipeline quickly locally. To facilitate this in the current project, there was a configuration flag that determined whether Dataflow and TensorFlow tasks were run locally or on the cloud, and whether the filesystem used was local or GCS. The functions in `tensorflow.python.lib.io.file_io` can be used to write generic functions that operate on both GCS and local files.
- When a pipeline stage would fail, it could be easily re-run from the Airflow web UI. Yet often re-running the stage would fail because the previous attempt had written/deleted some files to/from a place that the re-running stage would also attempt to access. That would happen after a delay when the stage does some work, so finding out about an obvious problem was time-consuming. A neat solution to this was writing a “checks” context manager that would take as input the task name as a string, and the context manager would execute pre- and postcondition checks for that task. That way it was possible to define for most tasks certain conditions, such as there should (not) be a given file in a given location before/after the execution of the task. This reduced the time to find out about a potential problem or bug, and resulted in more useful error messages.

References

- [1] Sanjeev Arora, Mikhail Khodak, Nikunj Saunshi, and Kiran Vodrahalli. A compressed sensing view of unsupervised text embeddings, bag-of-n-grams, and LSTMs. In *International Conference on Learning Representations*, 2018.
- [2] Richard G. Baraniuk. Compressive sensing. In *42nd Annual Conference on Information Sciences and Systems, CISS 2008, Princeton, NJ, USA, 19-21 March 2008*, 2008.
- [3] Luca Bertinetto, Jack Valmadre, João F. Henriques, Andrea Vedaldi, and Philip H. S. Torr. Fully-convolutional siamese networks for object tracking. *CoRR*, abs/1606.09549, 2016.
- [4] Leonid Boytsov and Bilegsaikhan Naidan. Engineering efficient and effective non-metric space library. In *Similarity Search and Applications - 6th International Conference, SISAP 2013, A Coruña, Spain, October 2-4, 2013, Proceedings*, pages 280–293, 2013.
- [5] Leo Breiman. Bagging predictors. *Machine learning*, 24(2):123–140, 1996.
- [6] Eric Brochu, Vlad M. Cora, and Nando de Freitas. A tutorial on bayesian optimization of expensive cost functions, with application to active user modeling and hierarchical reinforcement learning. *CoRR*, abs/1012.2599, 2010.
- [7] Emmanuel J. Candès, Justin Romberg, and Terence Tao. Robust uncertainty principles: Exact signal reconstruction from highly incomplete frequency information. *Information Theory, IEEE Transactions on*, 52(2):489–509, 2006.
- [8] Daniel Cer, Yinfei Yang, Sheng-yi Kong, Nan Hua, Nicole Limtiaco, Rhomni St. John, Noah Constant, Mario Guajardo-Cespedes, Steve Yuan, Chris Tar, Yun-Hsuan Sung, Brian Strope, and Ray Kurzweil. Universal sentence encoder. *CoRR*, abs/1803.11175, 2018.
- [9] Heng-Tze Cheng, Levent Koc, Jeremiah Harmsen, Tal Shaked, Tushar Chandra, Hrishi Aradhye, Glen Anderson, Greg Corrado, Wei Chai, Mustafa Ispir, Rohan Anil, Zakaria Haque, Lichan Hong, Vihan Jain, Xiaobing Liu, and Hemal Shah. Wide & deep learning for recommender systems. *CoRR*, abs/1606.07792, 2016.

REFERENCES

- [10] Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- [11] François Chollet. Information-theoretical label embeddings for large-scale image classification. *CoRR*, abs/1607.05691, 2016.
- [12] Ido Dagan and Sean P Engelson. Committee-based sampling for training probabilistic classifiers. In *Machine Learning Proceedings 1995*, pages 150–157. Elsevier, 1995.
- [13] Jesse Davis and Mark Goadrich. The relationship between precision-recall and roc curves. In *Proceedings of the 23rd International Conference on Machine Learning, ICML '06*, pages 233–240, New York, NY, USA, 2006. ACM.
- [14] Ali Mamdouh Elkahky, Yang Song, and Xiaodong He. A multi-view deep learning approach for cross domain user modeling in recommendation systems. In *Proceedings of the 24th International Conference on World Wide Web, WWW '15*, pages 278–288, Republic and Canton of Geneva, Switzerland, 2015. International World Wide Web Conferences Steering Committee.
- [15] Yoav Freund and Robert E Schapire. A decision-theoretic generalization of online learning and an application to boosting. *Journal of computer and system sciences*, 55(1):119–139, 1997.
- [16] Charlie Frogner, Chiyuan Zhang, Hossein Mobahi, Mauricio Araya-Polo, and Tomaso A. Poggio. Learning with a wasserstein loss. *CoRR*, abs/1506.05439, 2015.
- [17] Surya Ganguli and Haim Sompolinsky. Compressed sensing, sparsity, and dimensionality in neuronal information processing and data analysis. *Annual review of neuroscience*, 35:485–508, 2012.
- [18] Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge. A neural algorithm of artistic style. *CoRR*, abs/1508.06576, 2015.
- [19] Jonas Gehring, Michael Auli, David Grangier, Denis Yarats, and Yann N. Dauphin. Convolutional sequence to sequence learning. *CoRR*, abs/1705.03122, 2017.
- [20] Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT press Cambridge, 2016.
- [21] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In Z. Ghahramani, M. Welling, C. Cortes, N. D. Lawrence, and

REFERENCES

- K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems 27*, pages 2672–2680. Curran Associates, Inc., 2014.
- [22] Yves Grandvalet and Yoshua Bengio. Semi-supervised learning by entropy minimization. In L. K. Saul, Y. Weiss, and L. Bottou, editors, *Advances in Neural Information Processing Systems 17*, pages 529–536. MIT Press, 2005.
 - [23] Kyu J. Han, Akshay Chandrashekaran, Jungsuk Kim, and Ian R. Lane. The CAPIO 2017 conversational speech recognition system. *CoRR*, abs/1801.00059, 2018.
 - [24] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. *CoRR*, abs/1512.03385, 2015.
 - [25] Ruining He and Julian McAuley. Vbpr: Visual bayesian personalized ranking from implicit feedback. In *AAAI*, pages 144–150, 2016.
 - [26] Geoffrey E Hinton, Peter Dayan, Brendan J Frey, and Radford M Neal. The " wake-sleep" algorithm for unsupervised neural networks. *Science*, 268(5214):1158–1161, 1995.
 - [27] Geoffrey E Hinton, James L McClelland, David E Rumelhart, et al. Distributed representations. *Parallel distributed processing: Explorations in the microstructure of cognition*, 1(3):77–109, 1986.
 - [28] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
 - [29] Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, and Hartwig Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. *CoRR*, abs/1704.04861, 2017.
 - [30] Zhiting Hu, Zichao Yang, Xiaodan Liang, Ruslan Salakhutdinov, and Eric P. Xing. Toward controlled generation of text. In Doina Precup and Yee Whye Teh, editors, *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1587–1596, International Convention Centre, Sydney, Australia, 06–11 Aug 2017. PMLR.
 - [31] Gao Huang, Zhuang Liu, and Kilian Q. Weinberger. Densely connected convolutional networks. *CoRR*, abs/1608.06993, 2016.
 - [32] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *arXiv preprint arXiv:1502.03167*, 2015.

REFERENCES

- [33] Mohit Iyyer, Varun Manjunatha, Jordan Boyd-Graber, and Hal Daumé III. Deep unordered composition rivals syntactic methods for text classification. In *Association for Computational Linguistics*, 2015.
- [34] Eric Jang, Shixiang Gu, and Ben Poole. Categorical reparameterization with gumbel-softmax. *arXiv preprint arXiv:1611.01144*, 2016.
- [35] Surya Mattu Jeff Larson. How We Analyzed the COMPAS Recidivism Algorithm. <https://www.propublica.org/article/how-we-analyzed-the-compas-recidivism-algorithm>, 2016. [Online; accessed 17-May-2018].
- [36] Geoffrey E. Hinton Jimmy Lei Ba, Jamie Ryan Kiros. Layer normalization. 2016.
- [37] Brian Keng. Semi-supervised Learning with Variational Autoencoders. <http://bjlkeng.github.io/posts/semi-supervised-learning-with-variational-autoencoders/>, 2017. [Online; accessed 1-April-2018].
- [38] Murphy Kevin. Machine learning: a probabilistic perspective, 2012.
- [39] Yoon Kim. Convolutional neural networks for sentence classification. *arXiv preprint arXiv:1408.5882*, 2014.
- [40] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [41] Diederik P Kingma, Shakir Mohamed, Danilo Jimenez Rezende, and Max Welling. Semi-supervised learning with deep generative models. In *Advances in Neural Information Processing Systems*, pages 3581–3589, 2014.
- [42] Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- [43] Simon Kornblith, Jonathon Shlens, and Quoc V. Le. Do better imagenet models transfer better? 2018.
- [44] Ashutosh Kumar, Arijit Biswas, and Subhajit Sanyal. ecommercegan : A generative adversarial network for e-commerce. *CoRR*, abs/1801.03244, 2018.
- [45] Quoc V. Le and Tomas Mikolov. Distributed representations of sentences and documents. *CoRR*, abs/1405.4053, 2014.
- [46] Ilya Loshchilov and Frank Hutter. Fixing weight decay regularization in adam. *CoRR*, abs/1711.05101, 2017.
- [47] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. *CoRR*, abs/1301.3781, 2013.

REFERENCES

- [48] Matej Moravčík, Martin Schmid, Neil Burch, Viliam Lisý, Dustin Morrill, Nolan Bard, Trevor Davis, Kevin Waugh, Michael Johanson, and Michael H. Bowling. Deepstack: Expert-level artificial intelligence in no-limit poker. *CoRR*, abs/1701.01724, 2017.
- [49] Christopher Olah. Understanding LSTM Networks. <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>, 2015. [Online; accessed 30-March-2018].
- [50] Razvan Pascanu, Tomas Mikolov, and Yoshua Bengio. Understanding the exploding gradient problem. *CoRR*, abs/1211.5063, 2012.
- [51] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. Glove: Global vectors for word representation. In *Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543, 2014.
- [52] Alec Radford, Luke Metz, and Soumith Chintala. Unsupervised representation learning with deep convolutional generative adversarial networks. *arXiv preprint arXiv:1511.06434*, 2015.
- [53] Joseph Redmon and Ali Farhadi. YOLO9000: better, faster, stronger. *CoRR*, abs/1612.08242, 2016.
- [54] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, Alexander C. Berg, and Li Fei-Fei. Imagenet large scale visual recognition challenge. *Int. J. Comput. Vision*, 115(3):211–252, December 2015.
- [55] Jan Rygl, Jan Pomikálek, Radim Rehurek, Michal Ruzicka, Vít Novotný, and Petr Sojka. Semantic vector encoding and similarity search using fulltext search engines. *CoRR*, abs/1706.00957, 2017.
- [56] M. L. Seltzer and J. Droppo. Multi-task learning in deep neural networks for improved phoneme recognition. In *2013 IEEE International Conference on Acoustics, Speech and Signal Processing*, pages 6965–6969, May 2013.
- [57] Burr Settles. Active learning literature survey. Technical report, 2010.
- [58] Burr Settles, Mark Craven, and Soumya Ray. Multiple-instance active learning. In *Advances in neural information processing systems*, pages 1289–1296, 2008.
- [59] David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel, and Demis Hassabis. Mastering the game of go without human knowledge. *Nature*, 550:354–, October 2017.

REFERENCES

- [60] Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. *The Journal of Machine Learning Research*, 15(1):1929–1958, 2014.
- [61] Harini Suresh. Vanishing Gradients & LSTMs. <http://harinisuresh.com/2016/10/09/lstms/>, 2016. [Online; accessed 30-March-2018].
- [62] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott E. Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going deeper with convolutions. *CoRR*, abs/1409.4842, 2014.
- [63] Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jonathon Shlens, and Zbigniew Wojna. Rethinking the inception architecture for computer vision. *CoRR*, abs/1512.00567, 2015.
- [64] Antti Tarvainen and Harri Valpola. Weight-averaged consistency targets improve semi-supervised deep learning results. *CoRR*, abs/1703.01780, 2017.
- [65] Aäron van den Oord, Sander Dieleman, Heiga Zen, Karen Simonyan, Oriol Vinyals, Alex Graves, Nal Kalchbrenner, Andrew W. Senior, and Koray Kavukcuoglu. Wavenet: A generative model for raw audio. *CoRR*, abs/1609.03499, 2016.
- [66] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *CoRR*, abs/1706.03762, 2017.
- [67] Hao Wang, Naiyan Wang, and Dit-Yan Yeung. Collaborative deep learning for recommender systems. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1235–1244. ACM, 2015.
- [68] Jiang Wang, Yang Song, Thomas Leung, Chuck Rosenberg, Jingbin Wang, James Philbin, Bo Chen, and Ying Wu. Learning fine-grained image similarity with deep ranking. In *Proceedings of the 2014 IEEE Conference on Computer Vision and Pattern Recognition, CVPR ’14*, pages 1386–1393, Washington, DC, USA, 2014. IEEE Computer Society.
- [69] Xiang Wei, Zixia Liu, Liqiang Wang, and Boqing Gong. Improving the improved training of wasserstein GANs. In *International Conference on Learning Representations*, 2018.
- [70] Zhizheng Wu, Cassia Valentini-Botinhao, Oliver Watts, and Simon King. Deep neural networks employing multi-task learning and stacked bottleneck features for speech synthesis. In *Acoustics, Speech and Signal Processing (ICASSP), 2015 IEEE International Conference on*, pages 4460–4464. IEEE, 2015.

REFERENCES

- [71] Zichao Yang, Zhiting Hu, Ruslan Salakhutdinov, and Taylor Berg-Kirkpatrick. Improved variational autoencoders for text modeling using dilated convolutions. *arXiv preprint arXiv:1702.08139*, 2017.
- [72] Adams Wei Yu, David Dohan, Quoc Le, Thang Luong, Rui Zhao, and Kai Chen. Fast and accurate reading comprehension by combining self-attention and convolution. In *International Conference on Learning Representations*, 2018.
- [73] Fisher Yu and Vladlen Koltun. Multi-scale context aggregation by dilated convolutions. *CoRR*, abs/1511.07122, 2015.
- [74] Lantao Yu, Weinan Zhang, Jun Wang, and Yong Yu. Seqgan: Sequence generative adversarial nets with policy gradient. *CoRR*, abs/1609.05473, 2016.
- [75] Muhammad Bilal Zafar, Isabel Valera, Manuel Gomez Rogriguez, and Krishna P. Gummadi. Fairness Constraints: Mechanisms for Fair Classification. In Aarti Singh and Jerry Zhu, editors, *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, volume 54 of *Proceedings of Machine Learning Research*, pages 962–970, Fort Lauderdale, FL, USA, 20–22 Apr 2017. PMLR.

Declaration

I hereby certify that I have written this thesis independently and have only used the specified sources and resources indicated in the bibliography.

London, UK, July 2, 2018

.....
Mattias Arro

A. Additional Background

A.1 Neural Models

A.1.1 1D Convolutional Neural Networks

A simple convolutional architecture for text classification is described in [39]. Pre-trained k -dimensional word vectors are concatenated to represent a sentence, and a filter is $w \in R^{hk}$ is applied to a window of h words at each possible window of words in a sentence; this gives a single variable-length feature vector representing the sentence. Several such filters are learned (each with potentially a different width h), and max-over-time pooling is applied to these feature maps; this gives a vector of the highest activations from each filter, which is then passed to a fully-connected softmax or sigmoid layer for final classification.

This simple architecture should be able to predict output classes reasonably. Each filter can indicate the presence of a sequence of h words; max-pooling discards information about where exactly in the text it appeared, and ensures the output is of a fixed length. The same filter w is used across all possible word windows in the sentence, which can be seen as a form of parameter sharing (and as an infinitely strong prior over the parameters of the model [20]), and enables processing variable-length sequences. The relatively small number of shared parameters requires less training data than an equivalent fully-connected architecture. Using pre-trained word embeddings further simplifies the task, as these already carry some information about the meaning or syntactic role of words.

More sophisticated architectures can be built using 1D CNNs, although not used in our experiments. Most notably, several layers of convolutions and pooling can be stacked, where the following layer works on the feature maps output by the previous layer. The model described above uses a stride and dilation of 1, but multi-layer architectures where the higher layers have exponentially increasing dilation are able to expand the receptive field of the model, and as a result aggregate “contextual information from multiple scales” [73]. Dilated convolutions were beneficial for semantic segmentation of images with 2D CNNs; increased receptive field has also been beneficial in sequence-to-sequence NLP models [71].

A.1.2 Recurrent Neural Networks

Recurrent neural networks (RNNs) are a family of models for processing sequential data. The sequence of inputs $x^{(1)} \dots x^{(t)}$ in our case are vectors of word embeddings, but there are other options for input representations (e.g. sequence of 1-hot encoded tokens, or vectors representing a multivariate time series at a given time step t). Vanilla RNNs maintain a hidden state vector h which contains some information about the sequence it has seen so far, and is calculated from the input at the current time step t , and the previous hidden state:

$$h^{(t)} = f(h^{(t-1)}, x^{(t)}) \quad (\text{A.1})$$

Vanilla RNNs suffer from the vanishing and exploding gradient problem in long sequences and are notoriously difficult to train. Probably the most widely used variant, the long short-term memory networks (LSTM) [28] overcome this limitation by augmenting the network with a memory cell c ; there is a learned gating mechanism that controls what part of (and the extent to which) the cell is forgotten, what gets persisted in the cell state, and what gets output at a given time step. An intuitive overview of the gating mechanisms is given by [49] and [61] describes well how this helps mitigate vanishing gradients. A simpler gating mechanism is offered by gated recurrent units (GRU) [10], which has fewer parameters and works well on simpler tasks.

We are interested in using RNNs as text encoders, though they are often used for sequence to sequence modelling, or for predicting an output at each time step. In the simplest case, the concatenation of the cell state c and the hidden state h at the last time step could be used as a representation of the sequence.

A.2 Unsupervised and Semi-supervised Learning

This section explores ways in which unsupervised and semi-supervised learning can learn good feature representations and improve label complexity. Generic semi-supervised methods such as self-training are not considered, as one of our goals is to learn good representations, but also because these methods can reinforce poor predictions or do not make full use of all available unlabelled data.

Some generic methods can still improve the representations learned by our models when applied to models that learn deep embeddings in order to make predictions. Entropy minimisation [22] can be incorporated into models trained with SGD to encourage confident predictions of each class; this encourages the deep model to learn predictions that are strong predictors of some class and avoid producing features that produce mixed predictions. More complex but very performant approach is Mean Teacher [64]. A student that gets a harder task (such as predicting from a noisy/adversarial example), and a teacher that gets an easier task (i.e. the teacher model is an ensemble, which is more accurate). The student is trained on the prediction of the teacher; the teacher's parameters are an exponential moving av-

erage of the student’s, updated after each minibatch. The teacher’s predictions are higher-quality than the student’s (and can be applied on unlabelled data), while the student tries to continuously learn to learn a predictor that is robust to noisy and adversarial examples.

A.2.1 Autoencoders & Variational Autoencoders

Deep autoencoders (AEs) can be used to learn low-dimensional representations of inputs. Many variants exist, but the general pattern is to have an “hourglass-structured” neural network where the first half of the model shrinks the input, and the second half of the network reconstructs it. Activations in the middle layer correspond to a dense embedding of the sample; the shrinking layers need to throw away some of the detail in the input, yet persists enough for the expanding layers to be able to reconstruct the original input with some fidelity. The embedding contains information that is mostly unique to the data point, while the parameters of the encoding and decoding layers ensure that the reconstructed input is realistic. It is common to add noise (Gaussian, dropout) to the input or the intermediate layers to increase resilience to noisy inputs and to prevent the autoencoder simply memorising each data point.

If the bottleneck layer is unconstrained, it will use a wide range of values to represent different inputs. We would like embedding for similar inputs to also be similar, which is not always the case for ordinary autoencoders. Variational autoencoders (VAEs) impose a prior distribution on the values of the embedding z , often a multivariate Gaussian distribution with a diagonal covariance matrix, and the model is regularised during training to respect this prior. The model is optimised to minimise the reconstruction loss and KL divergence between the model’s distribution of and the prior on it, which can be computed just from $\mu(z)$ and $\sigma(z)$ if the prior is a isotropic standard normal. Enforcing the prior also means that the embeddings z will occupy a smooth, contiguous space, which allows us to draw samples from the prior $\epsilon \sim \mathcal{N}(0, 1)$ and use that as the input to the decoder - this gives us a generative model, which would not be possible with ordinary autoencoders. The output of the VAE is also probabilistic: e.g. for images, the output for a pixel would be a Gaussian distribution, which is sampled the same way as the Gaussians for z .

The VAEs described here are probabilistic models parameterised by neural networks for approximating the true posterior $p(z|x)$, since the exact posterior is intractable. In practice, VAEs add an extra loss term (KL divergence) to AEs, and additional step to determining the embedding z . There are two separate neural network layers from the bottleneck layer of the AE: one for determining $\mu(z)$ (the vector of means of the multivariate Gaussian), and one for determining $\sigma(z)$ (the vector of standard deviations of the multivariate gaussian). Given $\mu(z)$ and $\sigma(z)$, we can sample the final item embedding $z \sim \mathcal{N}(\mu, \sigma^2)$. Note that sampling of z is a discrete decision that would normally stop gradient from propagating past this step, but we can use the reparametrisation trick introduced by [42] that turns the

discrete decision into a deterministic function of $z = \mu + \sigma\epsilon$, where $\epsilon \sim \mathcal{N}(0, 1)$ is a random auxiliary noise parameter.

A.2.2 Conditional Variational Autoencoders

It is easier for linear models to learn from embeddings extracted with AEs, and embeddings from VAEs make the classes even easier to separate. These are unsupervised methods that do not take advantage of labels in the training data. In the semi-supervised approach introduced by [41] there are two inference networks: a discriminative classifier that outputs a categorical distribution from the input x , and a class-conditional encoder that takes as input both x and the 1-hot encoded categorical label y ¹. The model is trained on both labelled and unlabelled data. For labelled examples, the x and y are given as input to the model, which embeds it in z as described earlier, and reconstructs the original x . The value of y is unknown for unlabelled data points; therefore the model is run $|y|$ times, encoding and decoding the data point conditioned on each possible class i ; loss from these runs is averaged, weighted by the discriminator's estimate of the sample belonging to class i . This approach is acceptable for small numbers of classes, but is impractical for many multi-class problems. A solution uses the reparametrisation trick mentioned above: rather than running the procedure once per class, we take a Gumbel-softmax [34] to get a discrete estimate of y that is still differentiable.

Such class-conditional VAEs can be used in any situation where unlabelled data is abundant but labels are scarce. The discriminator's loss is optimised as part of the VAE's loss function, which means adding more unlabelled data can improve the accuracy of the discriminator. The approach was developed and tested on the MNIST dataset, which consists of 28x28 grayscale images - a very simple dataset by today's standards. Some reports have emerged that fail to reproduce this success on more complex datasets such as CIFAR-10, being outperformed even by PCA [37].

One related approach is offered by [30], where a class-conditional VAE is used to generate synthetic data for a discriminator network, which has reportedly a higher accuracy in semi-supervised setting than the approach described just above. They use it for conditional text generation, but this approach of using the class-conditional VAE to synthesise training data to train a discriminator is applicable to any input modality. In fact the general approach to “dreaming up” new training samples in the *sleep phase* and training the classifier in the *wake phase* was introduced already in [26]. The current author has implemented this model and open sourced it².

A.2.3 Generative Adversarial Networks

Generative adversarial networks (GANs) are an interesting approach capable of synthesising realistic data, but also useful for learning good representations of data.

¹Class-conditional encoder means that the encoder is aware of the class of the data point, i.e. the output distribution z is conditioned on the class y in addition to the original input x .

²<https://github.com/mattiasarro/seq2seq-cvae-tensorflow>

GANs have been most successful for image synthesis, but are in fact applicable to any kind of input data including text [74] and even mixed data types such as e-commerce orders [44]. In the GAN setting, there are two networks: a generator that tries to produce realistic synthetic samples, and a discriminator whose goal is to distinguish between synthetic and actual samples. Given that the training is stable enough, both the generator's and the discriminator's performance improves with time, resulting in more realistic synthetic samples, as opposed to the relative blurry images produced by VAEs. The discriminator needs to learn good representations of the input to accurately distinguish which samples come from the true distribution, and which samples are synthetic; as with many CNNs for computer vision, these representations can be useful for other tasks as well.

It has been shown that linear models from the embeddings learned by the discriminator outperform other unsupervised feature learning techniques such as k-means [52]. At the time of writing, variants of Wasserian GANs (WGANs) achieve the best performance by improving training stability and preventing the generator from generating samples from a limited number of modes; this is achieved by minimising the Wasserian distance between the generator's and real data distributions, as opposed to minimising the JS divergence as was common before. A performant variant of WGAN is CT-GAN, which adds a regularisation term to enforce a Lipschitz continuity condition over the manifold of the real data [69].